


Deep learning

Convolutional Neural
Networks

- Lecture


Basics of convolution
operation






Computer and Human Interpretation

Human
Interpretation



=



Computer
Interpretation



-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

=

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	1	-1	-1	-1
-1	-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1



Computer Interpretation

Pixel wise
Matching



-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	X	-1	-1	-1	-1	X	X	-1
-1	X	X	-1	-1	X	X	-1	-1
-1	-1	X	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	X	-1	-1
-1	-1	X	X	-1	-1	X	X	-1
-1	X	X	-1	-1	-1	-1	X	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

Decision



-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1




-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	1	-1	-1
-1	-1	-1	1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

Computers are Literal



Feature matching for symbol 'X'





Convolution Operation


1	-1	-1
-1	1	-1
-1	-1	1

$$1 \times 1 = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1




Piece Matching of Features





Piece Matching of Features





Convolution Operation

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature



Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$1 \times 1 = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1		



Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\begin{bmatrix} -1 \end{bmatrix} \times \begin{bmatrix} -1 \end{bmatrix} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	



Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{-1} \times \boxed{-1} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1



Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{-1} \times \boxed{-1} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1		



Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{1} \times \boxed{1} = 1$$

1	1	1
1	1	

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{-1} \times \boxed{-1} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	1





Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{-1} \times \boxed{-1} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	1
1		



Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{-1} \times \boxed{-1} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	1
1	1	





Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{1} \times \boxed{1} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	1
1	1	1



Convolution Operation

1	-1	-1
-1	1	-1
-1	-1	1

$$\boxed{1} \times \boxed{1} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	1
1	1	1





Feature Mapping

These are the network parameters to be learned.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

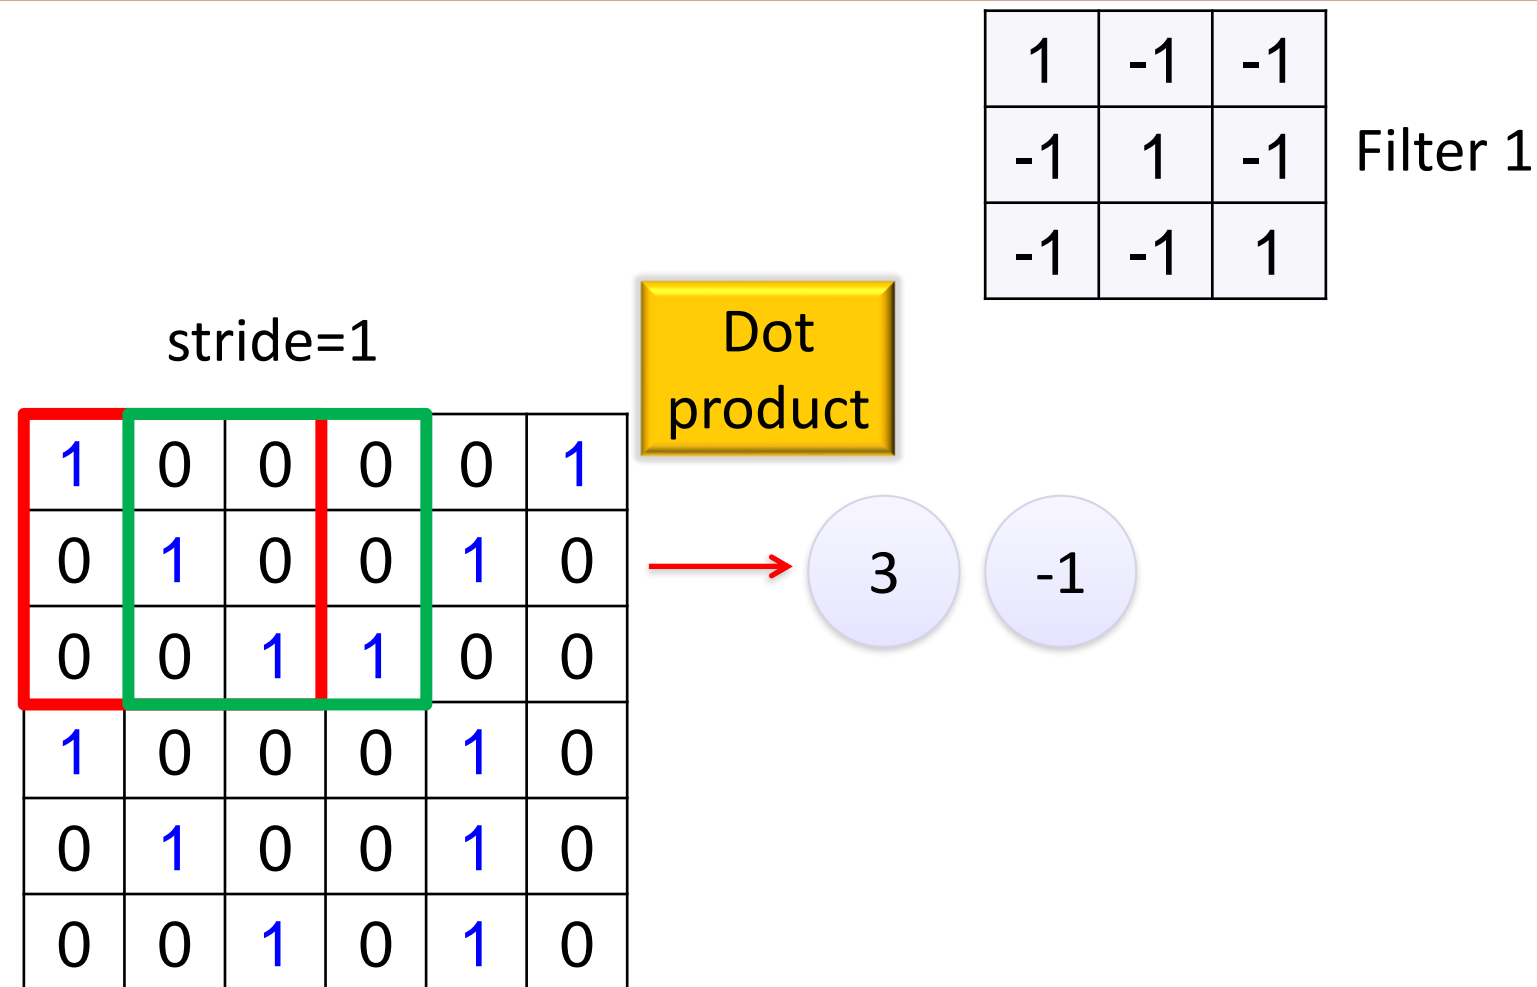
Filter 2

⋮

Each filter detects a small pattern (3 x 3).



Feature Mapping






Feature Mapping

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0



6 x 6 image



Feature Mapping

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1



Feature Mapping

stride=1


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2


Repeat this for each filter



Two 4 x 4 images
Forming 4 x 4 x 2 matrix





Convolution v.s. Fully Connected



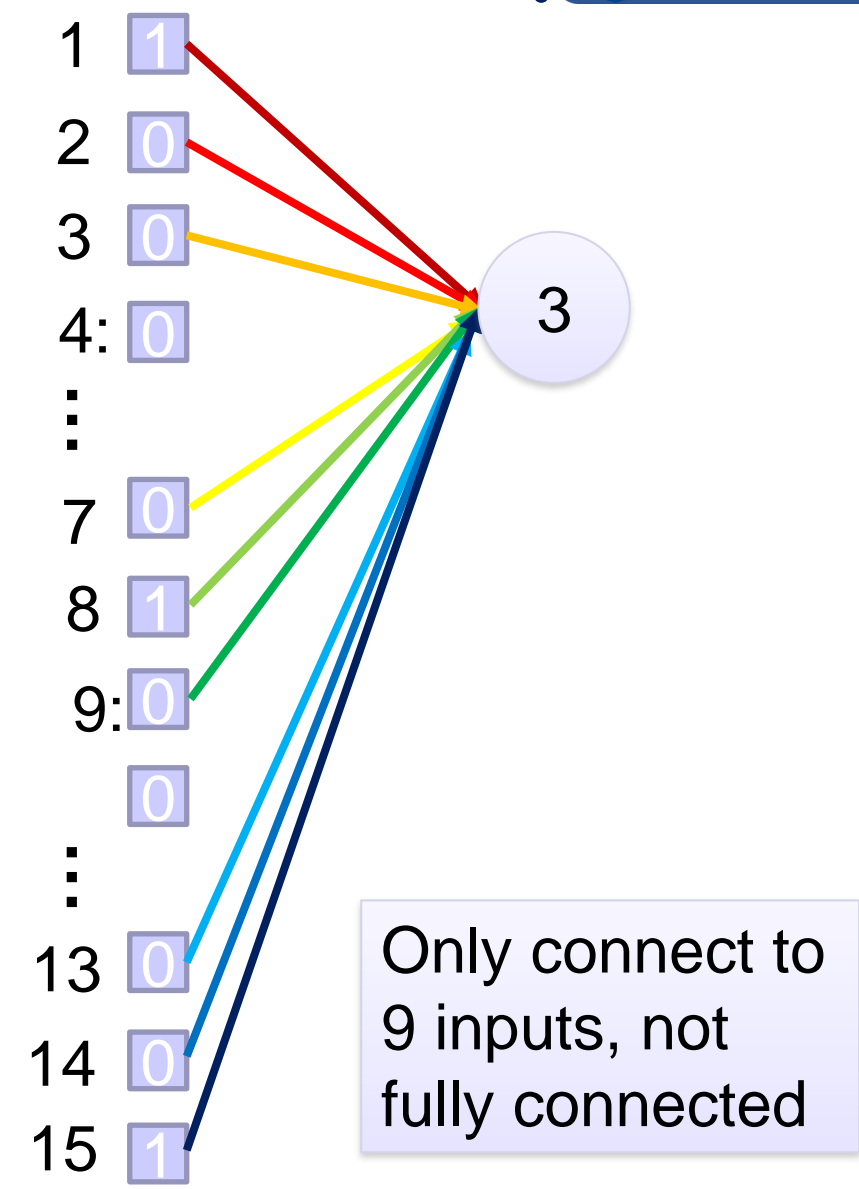
Fully-
connected



1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0





fewer parameters!





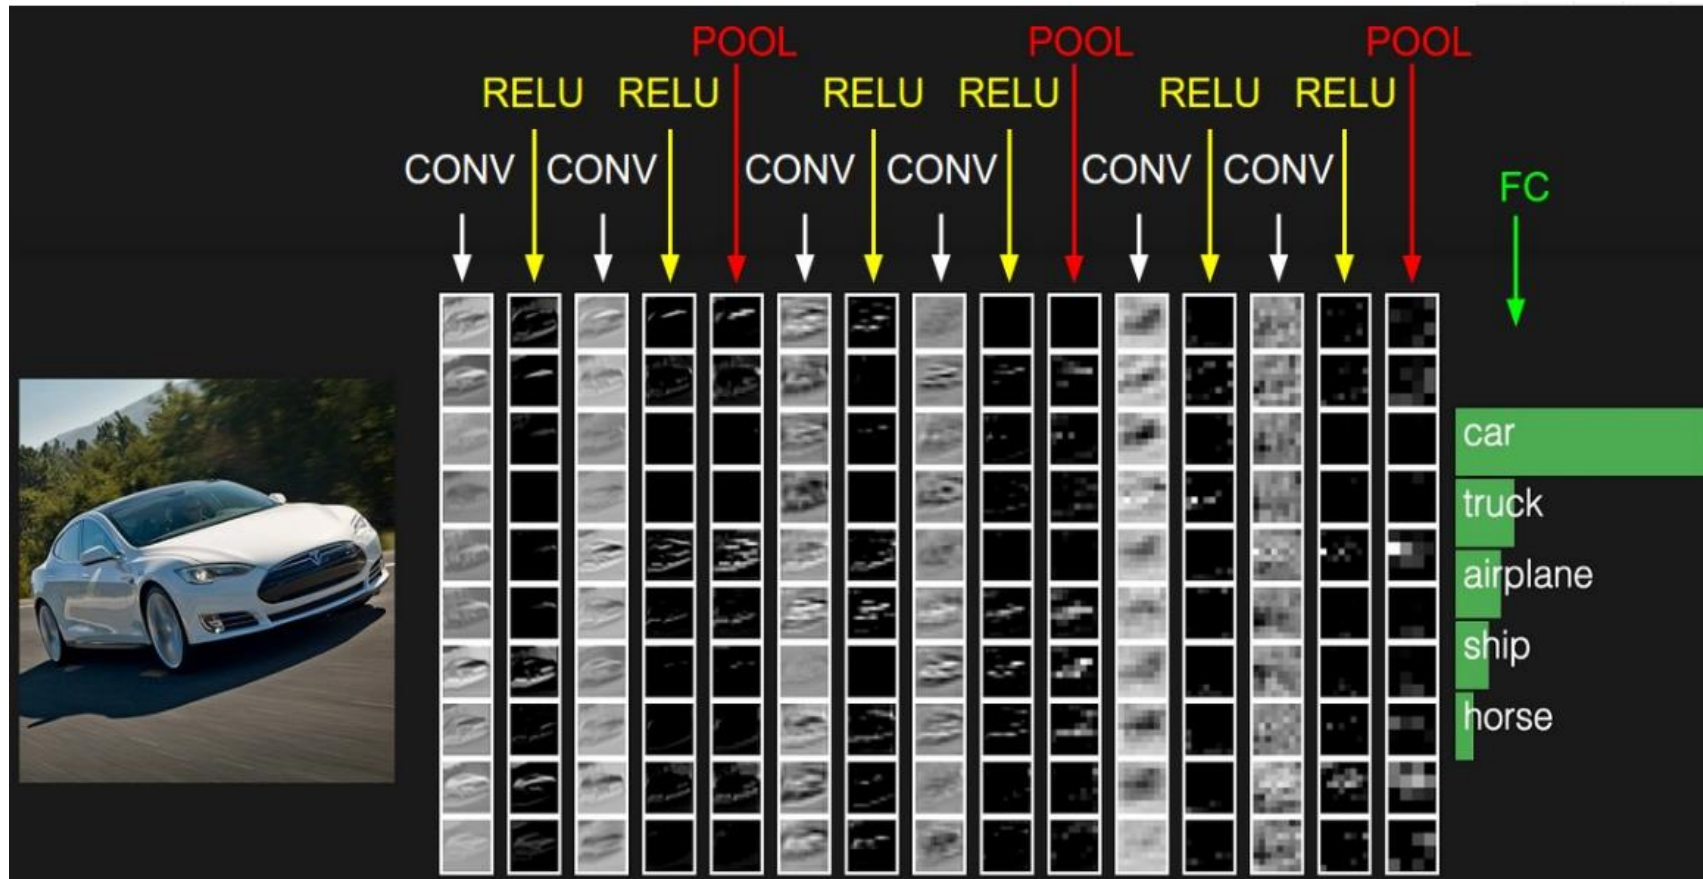
Shared weights

Fewer parameters

Even fewer parameters



Simple CNN architecture




CONV: Convolutional kernel layer

RELU: Activation function

POOL: Dimension reduction layer


FC: Fully connection layer

Convolutional Neural Network--- Spatial View



Stride = Step size of filter, Receptive Field = Location of connected path in an input image

Non-Linearity in Convolutional Neural Network

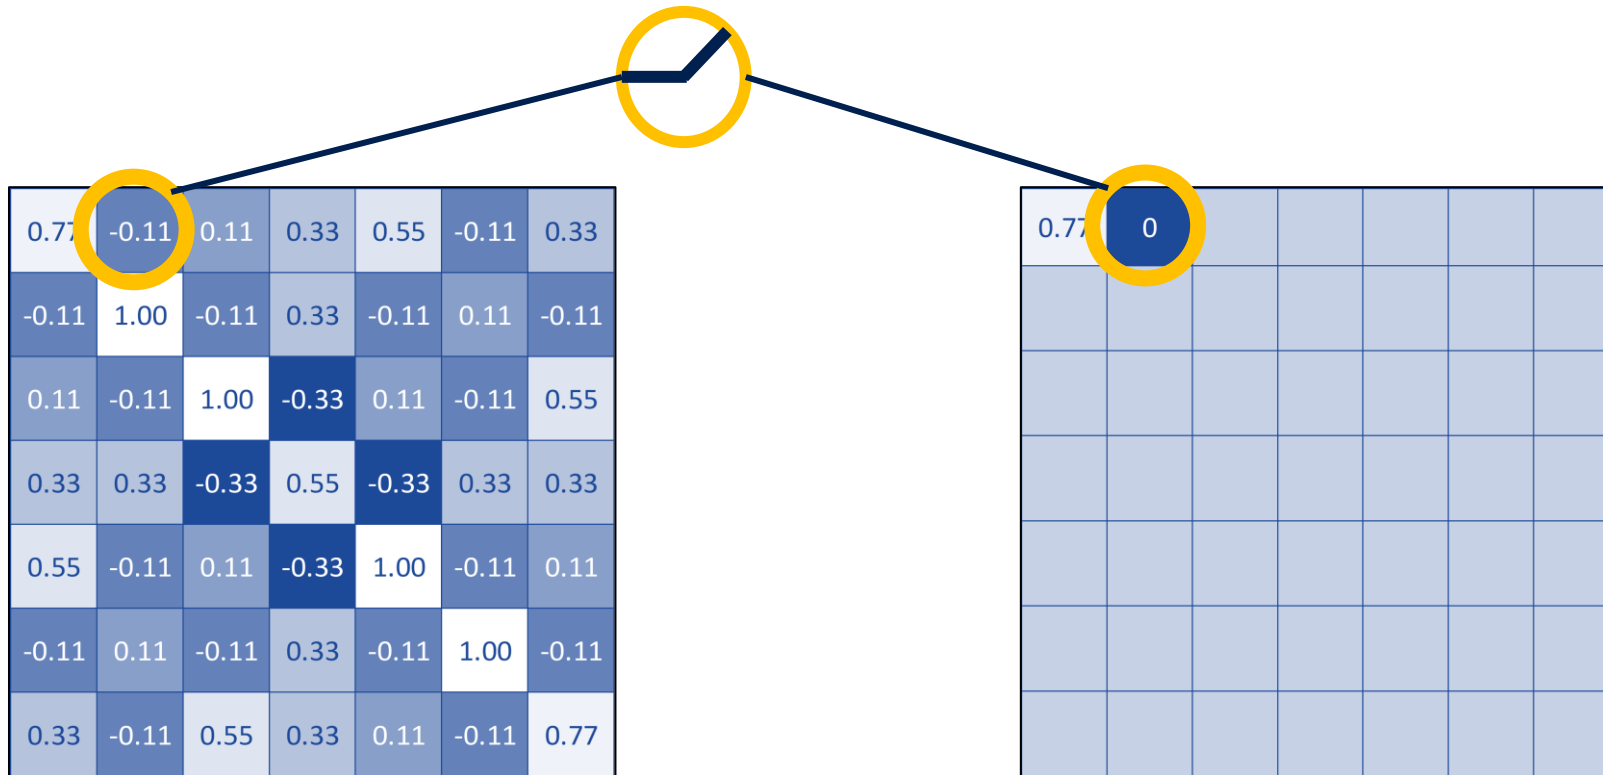


Applied after every convolutional layer. The rectified linear activation (ReLU) function is a simple calculation that returns the value provided as input directly, or the value 0.0 if the input is 0.0 or less.

$$g(x) = \max(0, x)$$




Rectified Linear Units (ReLUs)



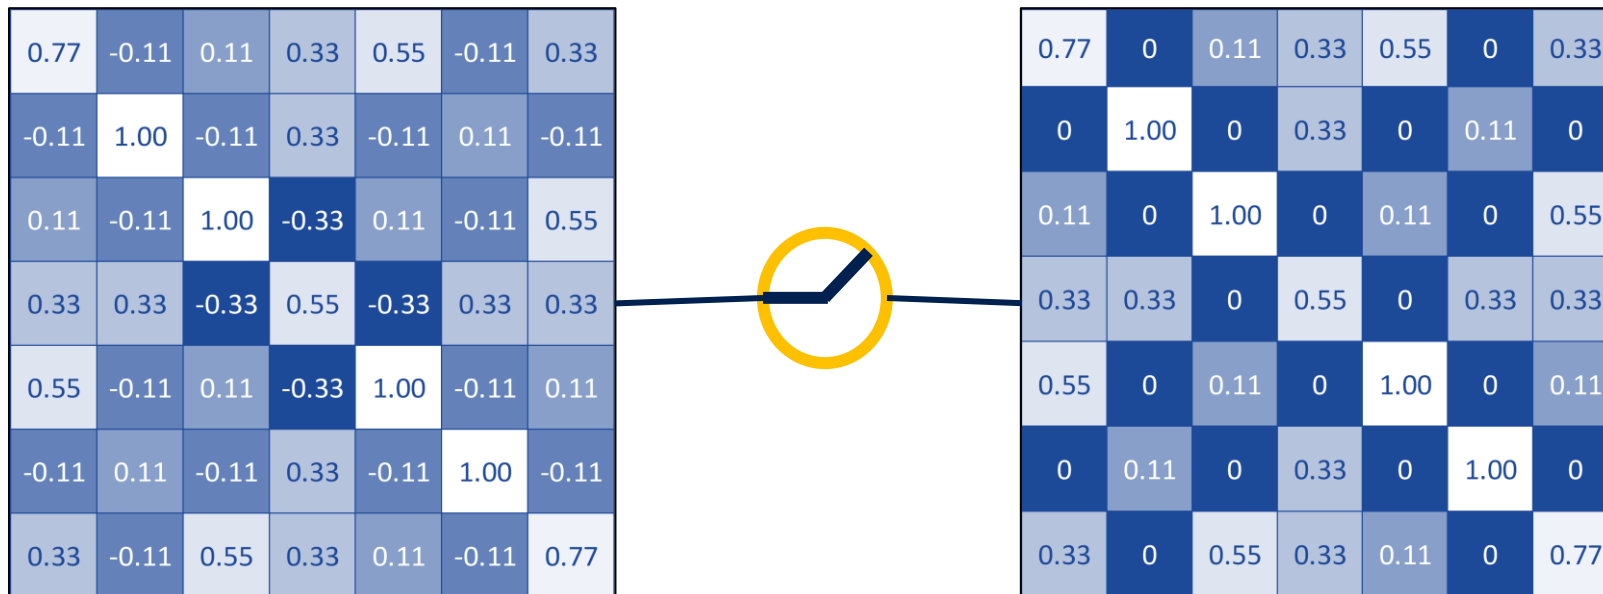


Rectified Linear Units (ReLUs)










Rectified Linear Units (ReLUs)





Rectified Linear Units (ReLUs)

A stack of images becomes a stack of images with no negative values.





Pooling

- *Dimensionality Reduction*
- *Preserve Spatial Invariance*

The types of pooling operations are:

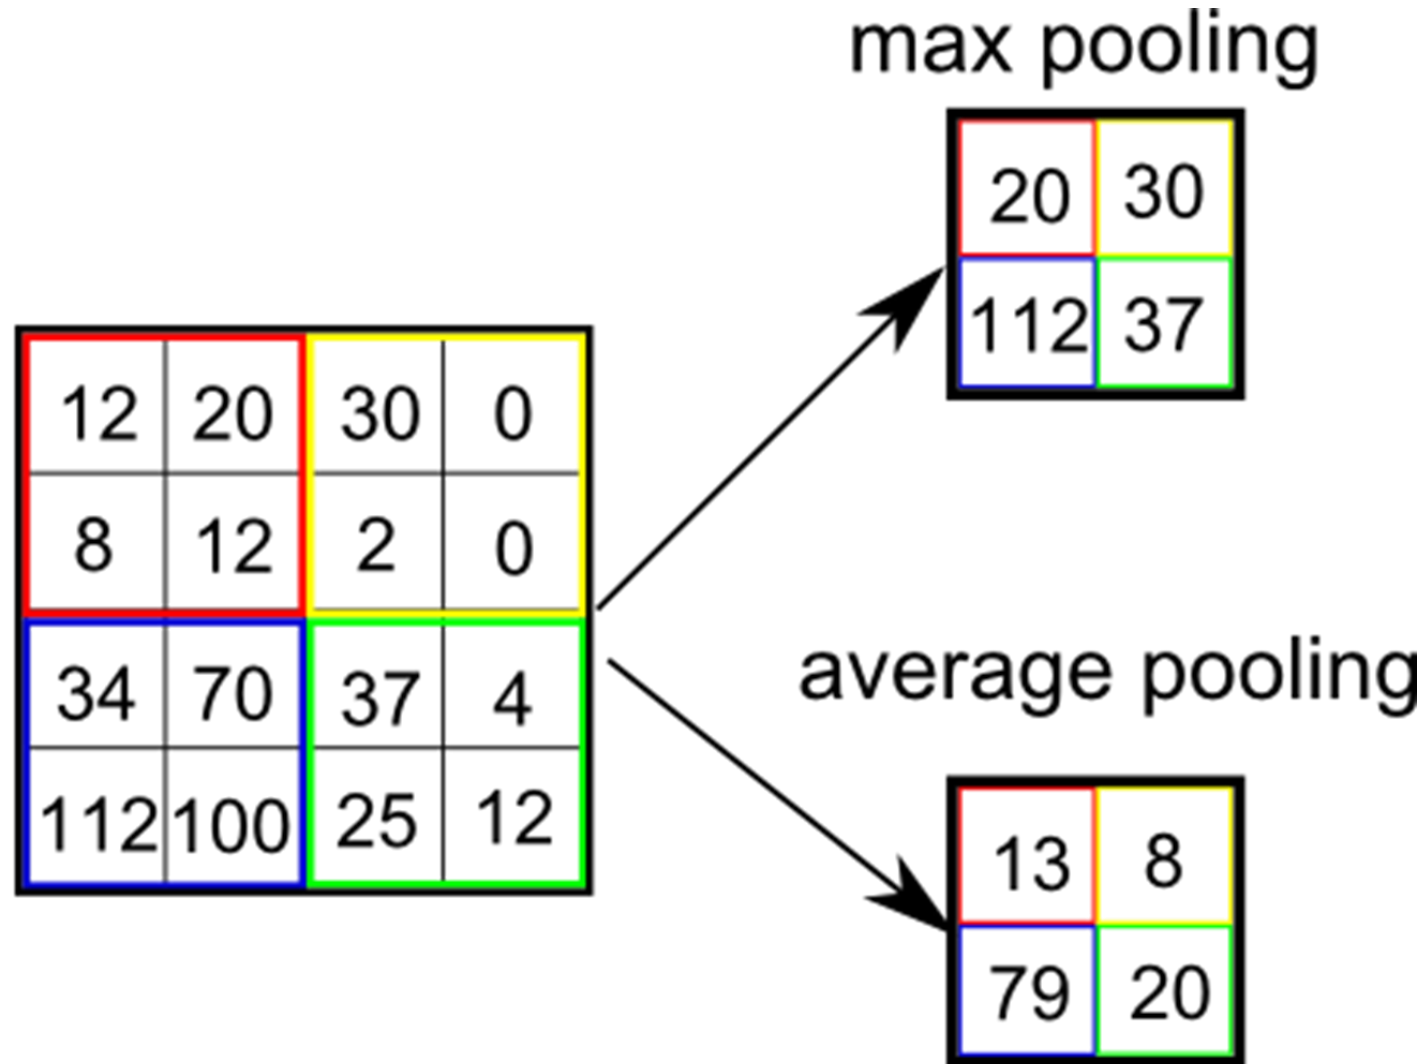
- ❑ **Max pooling:** The maximum pixel value of the batch is selected.
- ❑ **Average pooling:** The average value of all the pixels in the batch is selected.

STEPS

1. Pick a window size (usually 2 or 3).
2. Pick a stride (usually 2).
3. Walk your window across your filtered images.
4. From each window, take the maximum value.



Pooling





Why Pooling ?

- Subsampling pixels will not change the object

bird



Subsampling

bird

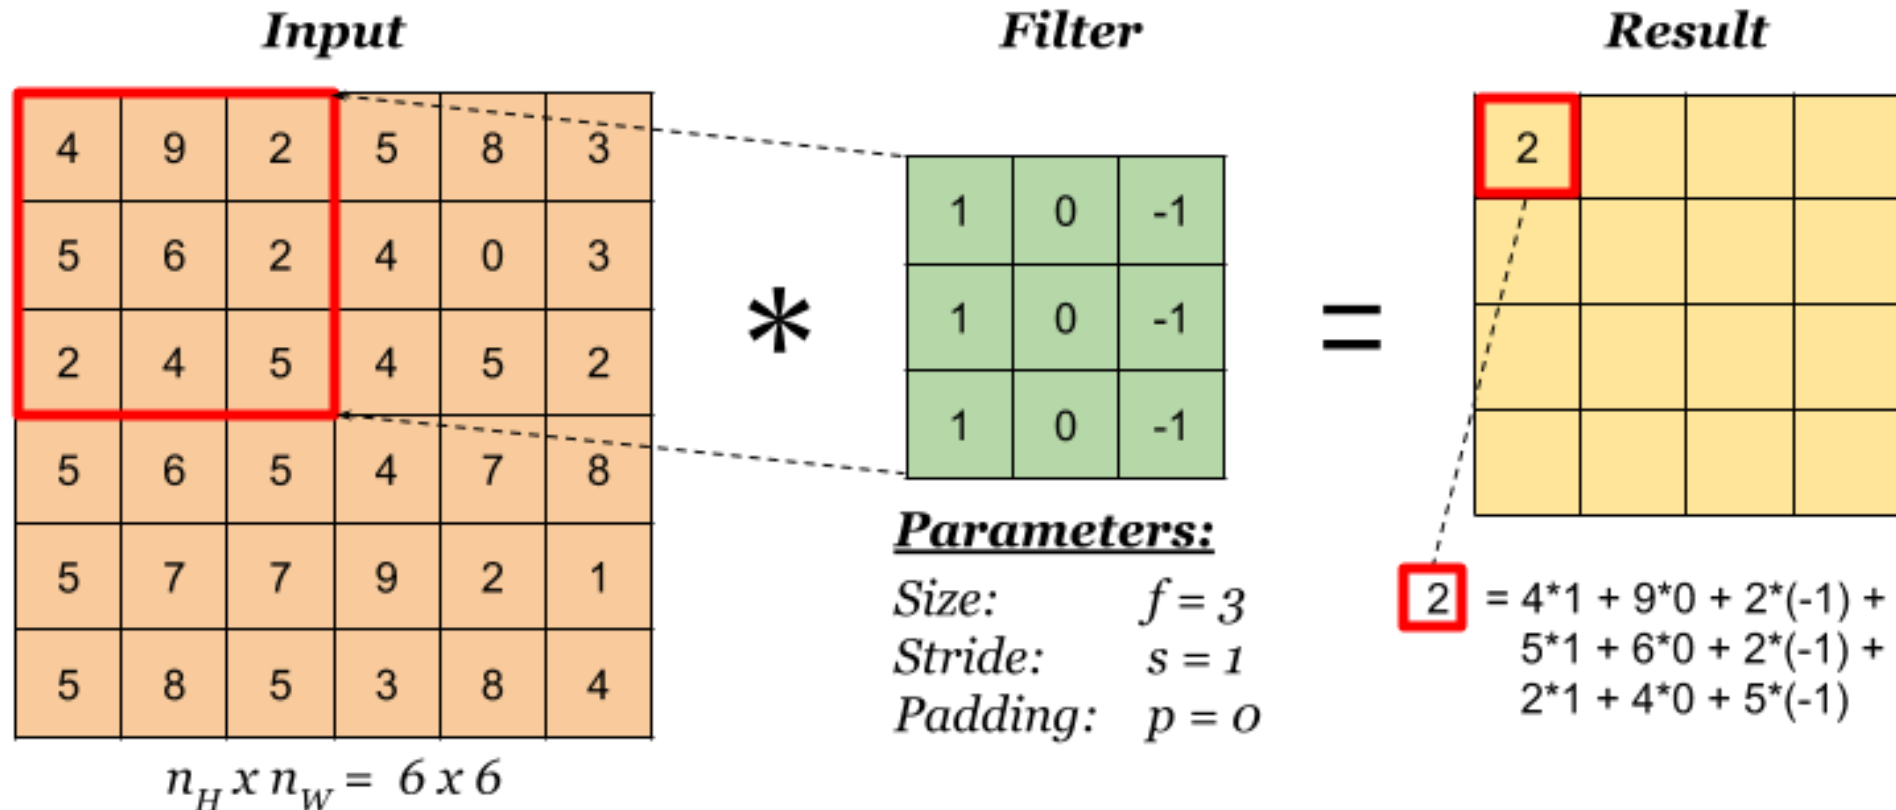


We can subsample the pixels to make image smaller fewer parameters to characterize the image



Stride

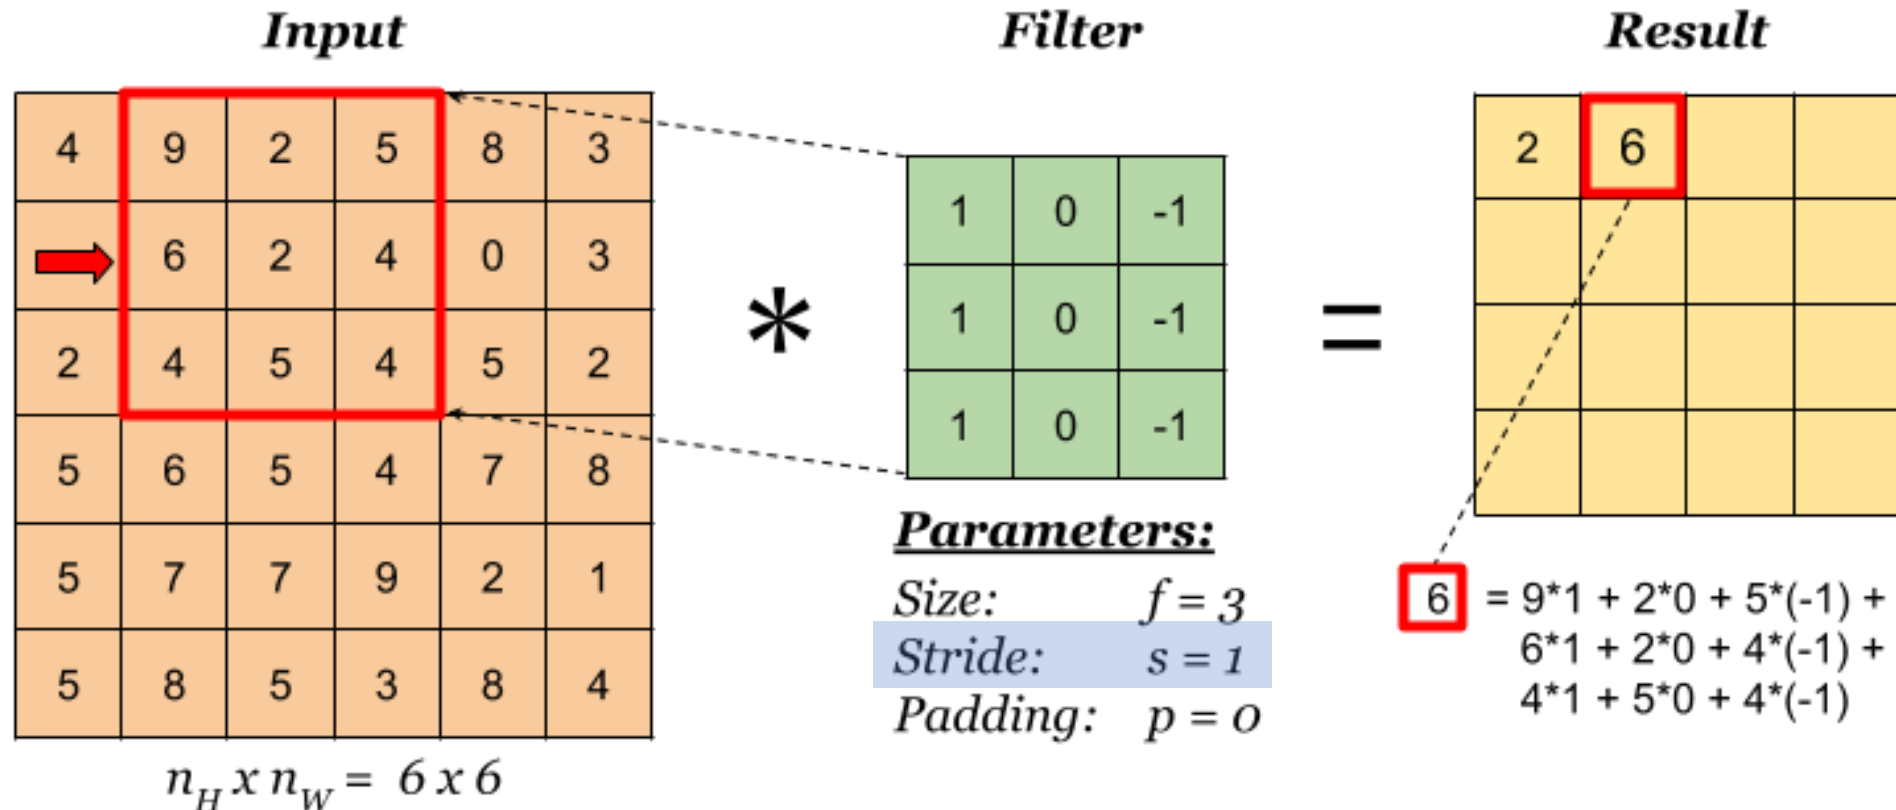
- number of cells the filter is moved to calculate the next output
- sample only every s pixels in each direction in the output





Stride

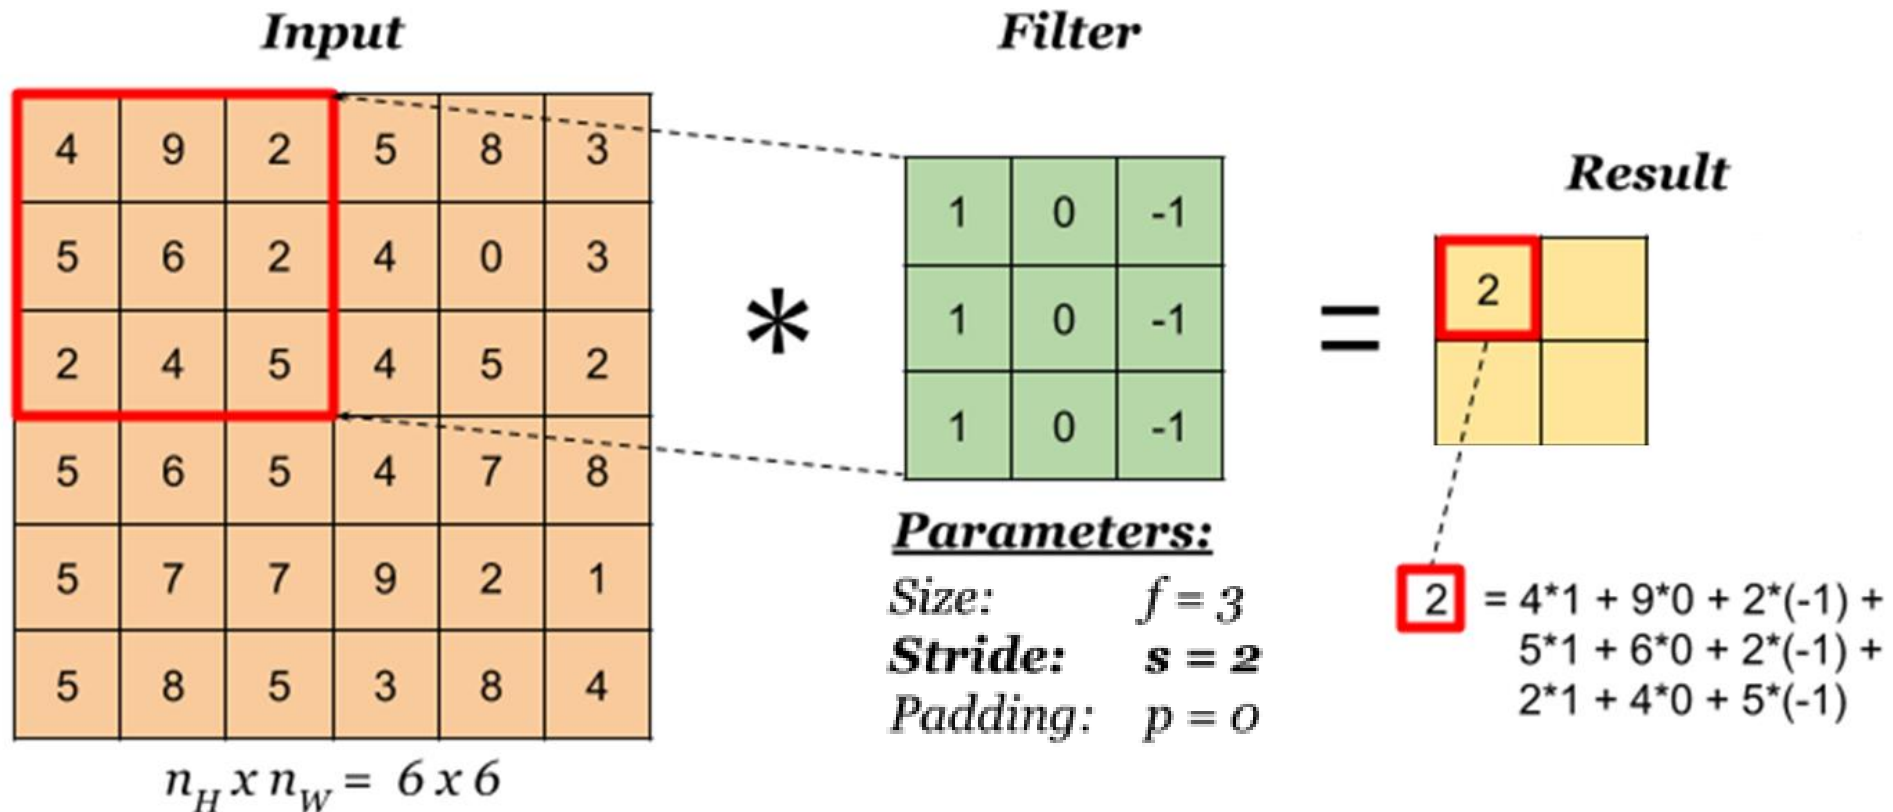
- number of cells the filter is moved to calculate the next output
- sample only every s pixels in each direction in the output





Stride

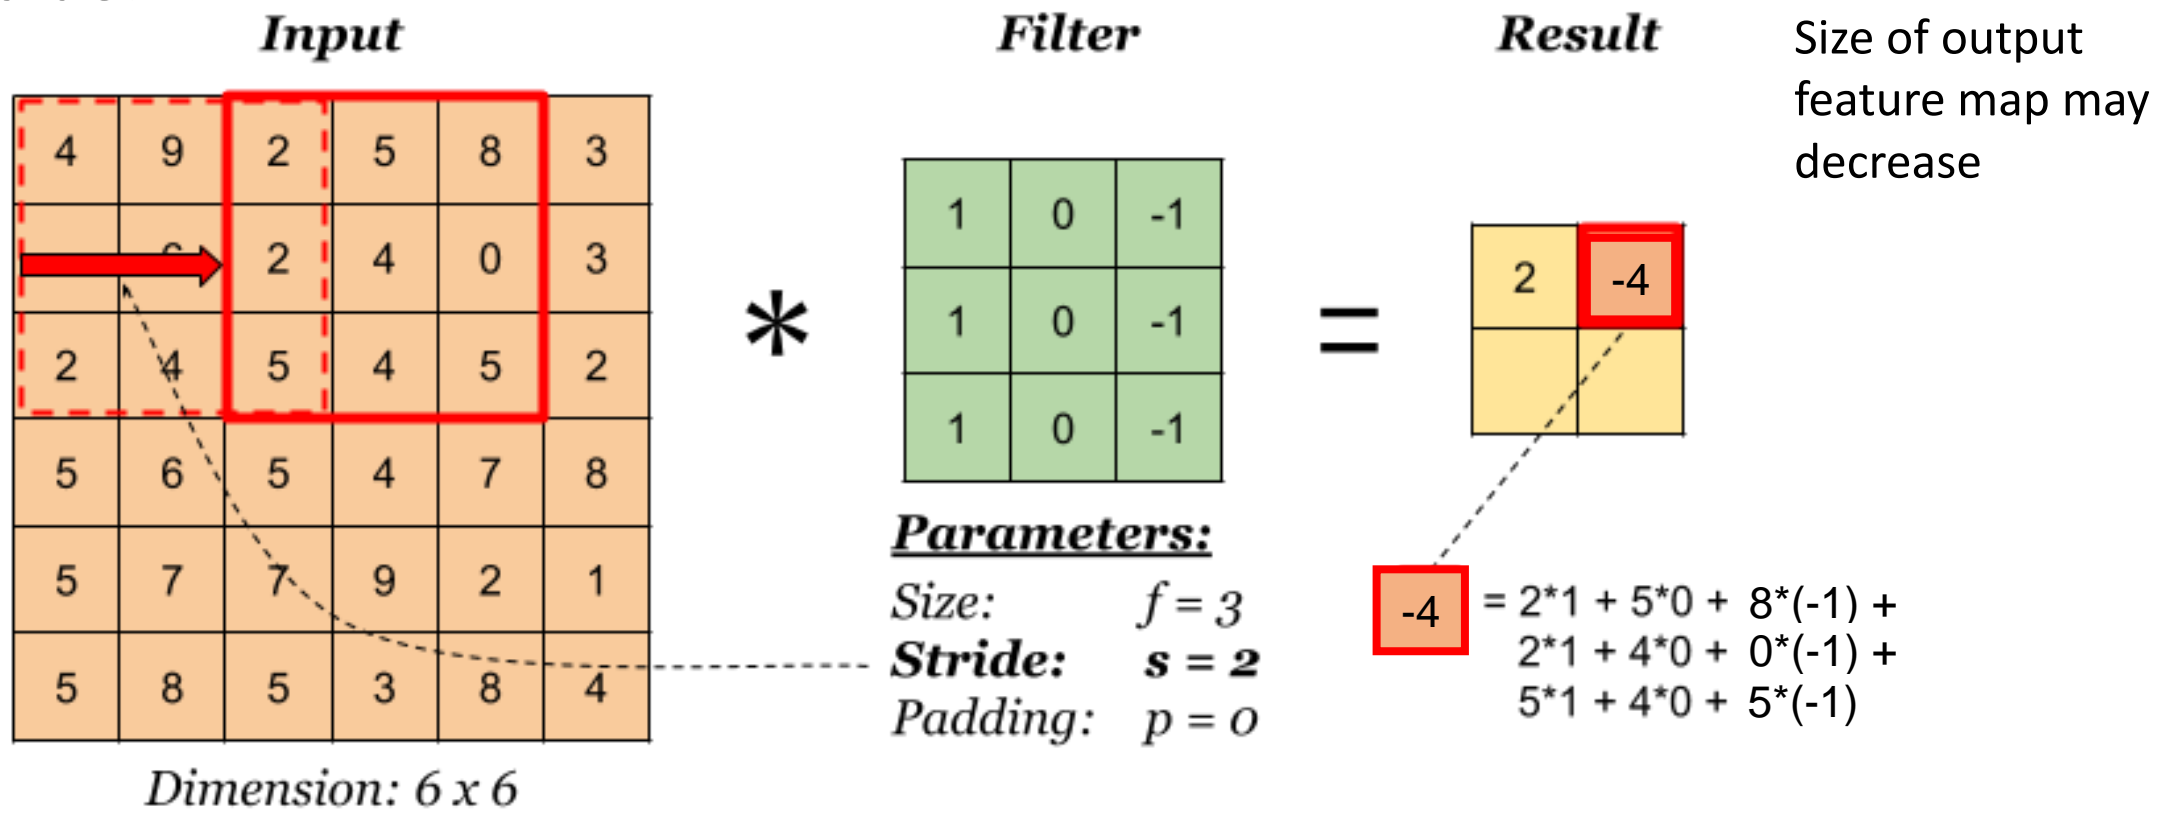
- Stride = 2
- First Value:





Stride

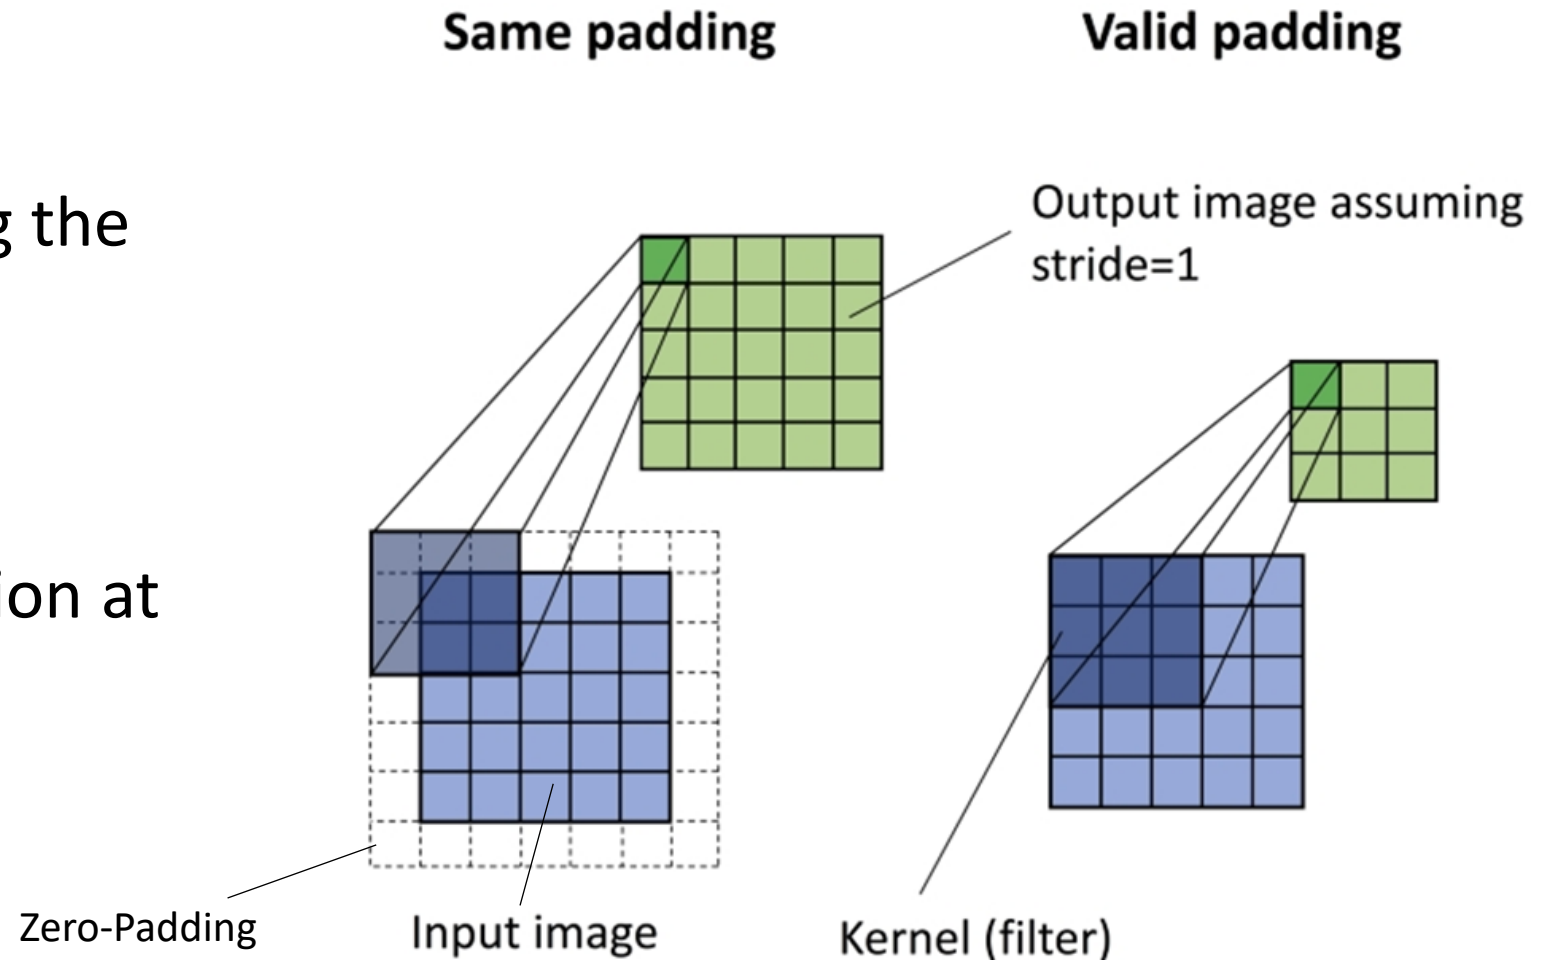
- Stride = 2
- Next Value:





Padding

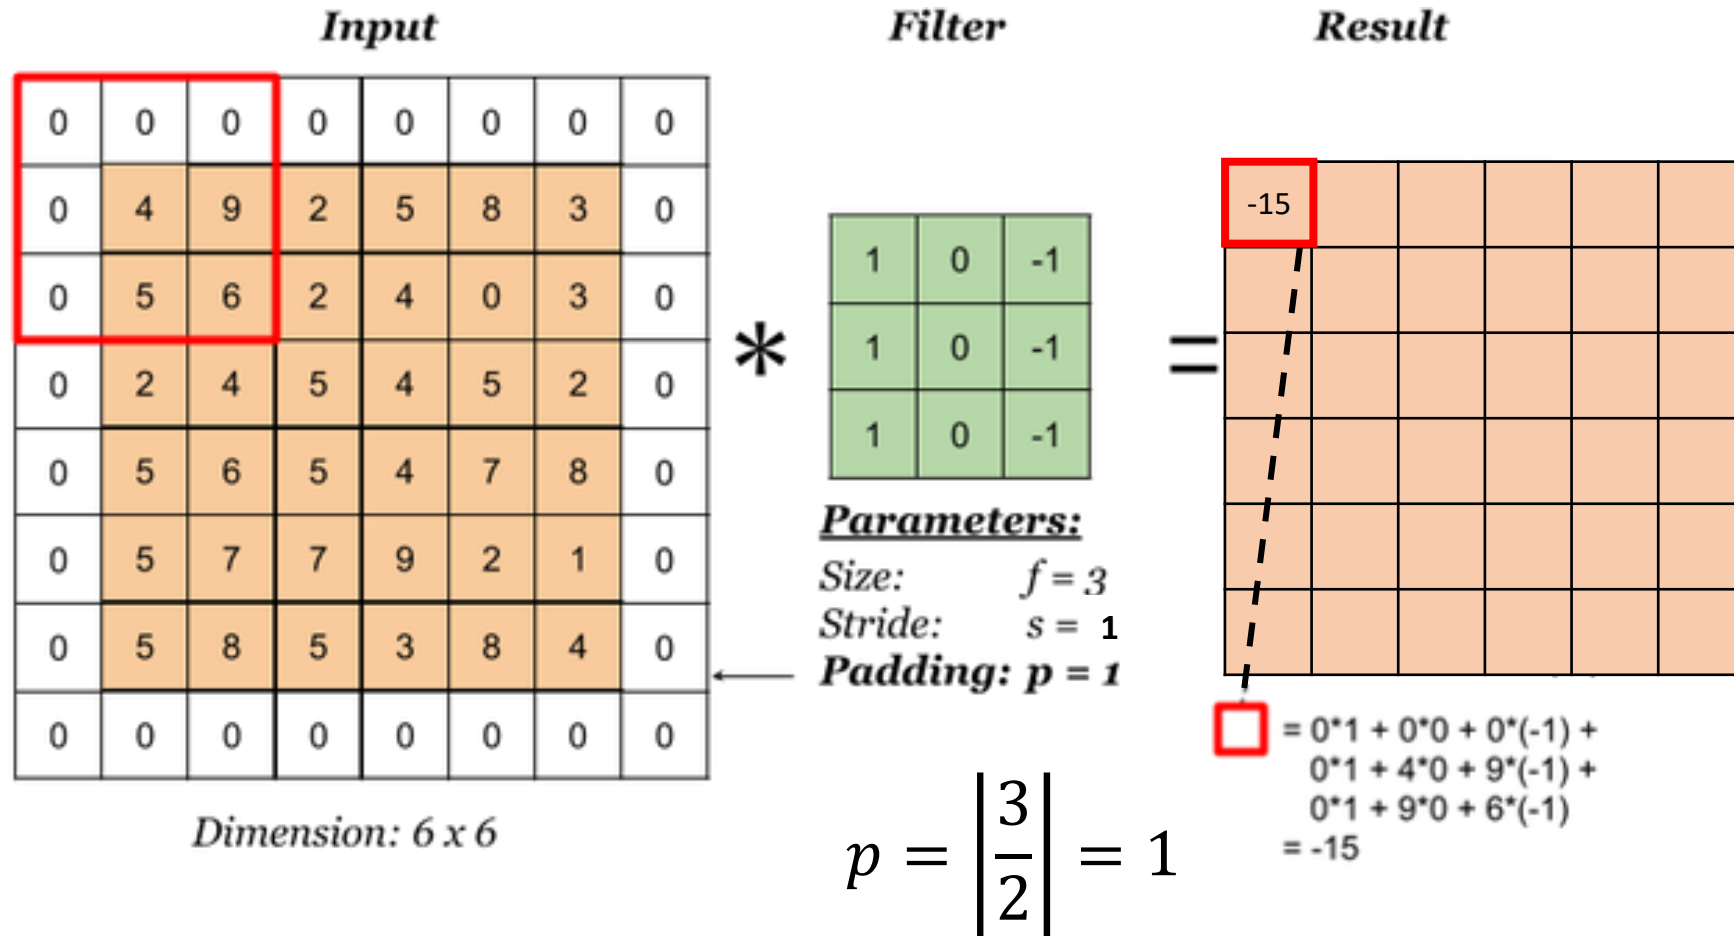
- Use Conv without shrinking the height and width
- Helpful in building deeper networks
- Keep more of the information at the border of an image





Same Padding


- Buffers the edge of the input with $\lfloor \text{filter_size}/2 \rfloor$ zeros (integer division)
- Output dimension is the same as the input for $s=1$





Same Padding

- Buffers the edge of the input with $\lfloor \text{filter_size}/2 \rfloor$ zeros (integer division)
- Output dimension is the same as the input for $s=1$
- Output dimension reduces less for $s>1$





HOW TO CALCULATE THE OUTPUT SIZE OF A CONVOLUTIONAL LAYER

To determine the output size of the convolution, the following equation can be applied:

$$n_{out} = \frac{n_{in} + 2p - k}{s} + 1$$



Putting it all together

```
import tensorflow as tf
```

```
def generate_model():
```

```
    model = tf.keras.Sequential([
```

```
        # first convolutional layer
```

```
        tf.keras.layers.Conv2D(32, filter_size=3, activation='relu'),
```

```
        tf.keras.layers.MaxPool2D(pool_size=2, strides=2),
```

```
        # second convolutional layer
```

```
        tf.keras.layers.Conv2D(64, filter_size=3, activation='relu'),
```

```
        tf.keras.layers.MaxPool2D(pool_size=2, strides=2),
```

```
        # fully connected classifier
```

```
        tf.keras.layers.Flatten(),
```


```
        tf.keras.layers.Dense(1024, activation='relu'),
```

```
        tf.keras.layers.Dense(10, activation='softmax')
```

```
        # 10 outputs
```

```
    ])
```

```
    return model
```



Deep learning

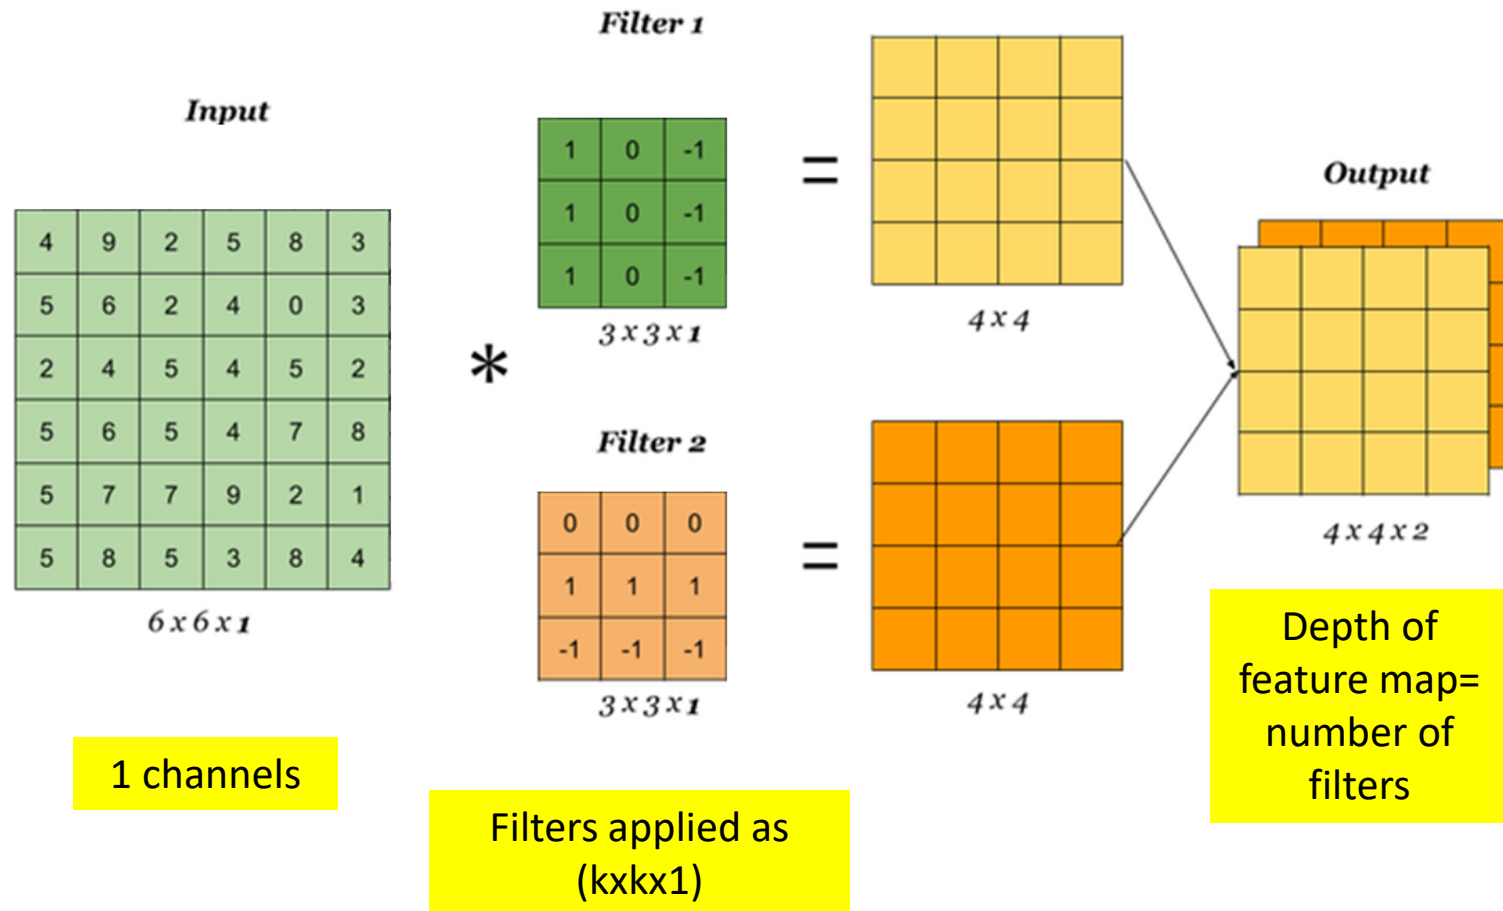
- Course Code:
- Unit 3
Convolutional Neural Networks
and Transfer Learning
- Lecture 4
Convolutional Neural Network





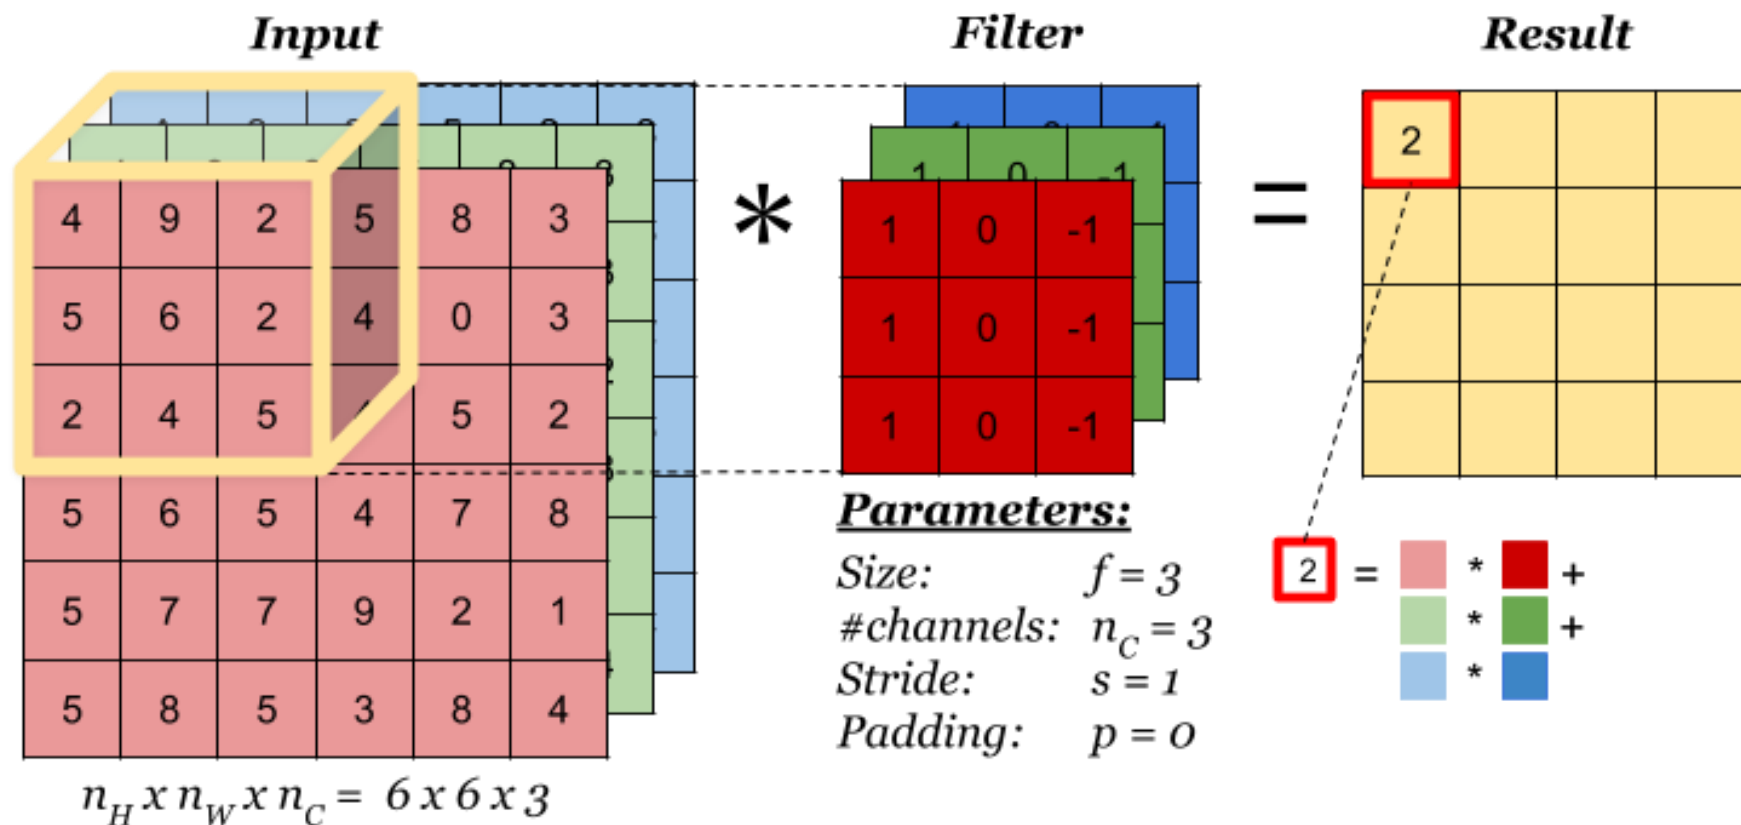
Convolution with Single Channel and Multiple Filters

- Input with 1 channel (eg. grayscale image) then a 3x3 filter will be applied in 3x3x1 blocks

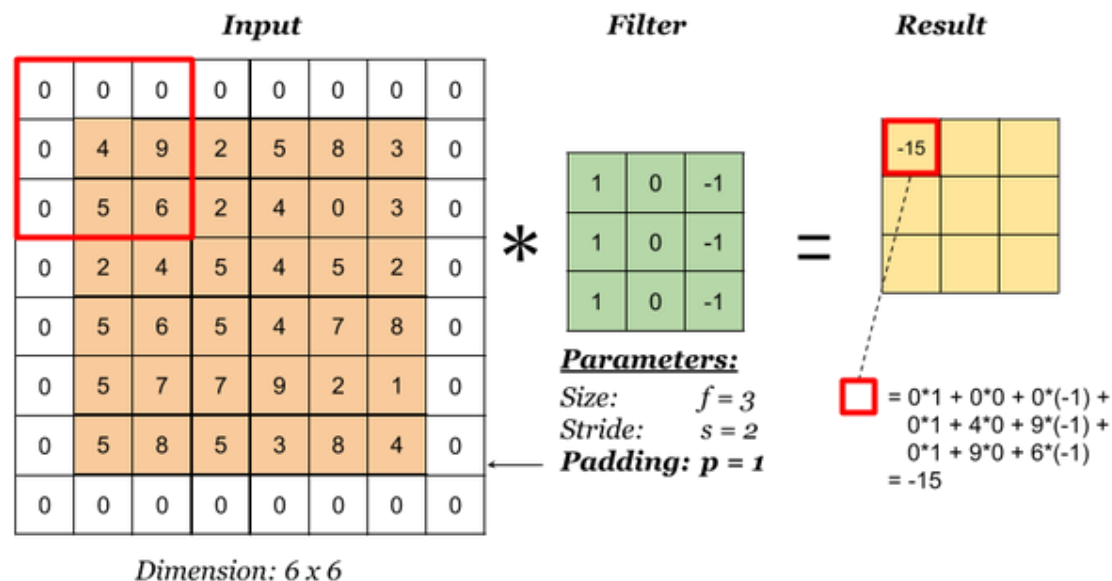


Convolution over Volume (Multiple Channels)

- RGB images has 3 channels:
Red, Green, Blue
- One kernel for every input channel to the layer (each kernel is unique)
- Each filter = a *collection of kernels*



Output Dimension



$$n_{out} = \left\lfloor \frac{n_{in} + 2p - k}{s} \right\rfloor + 1$$

n_{in} : number of input features

n_{out} : number of output features

k : convolution kernel size

p : convolution padding size

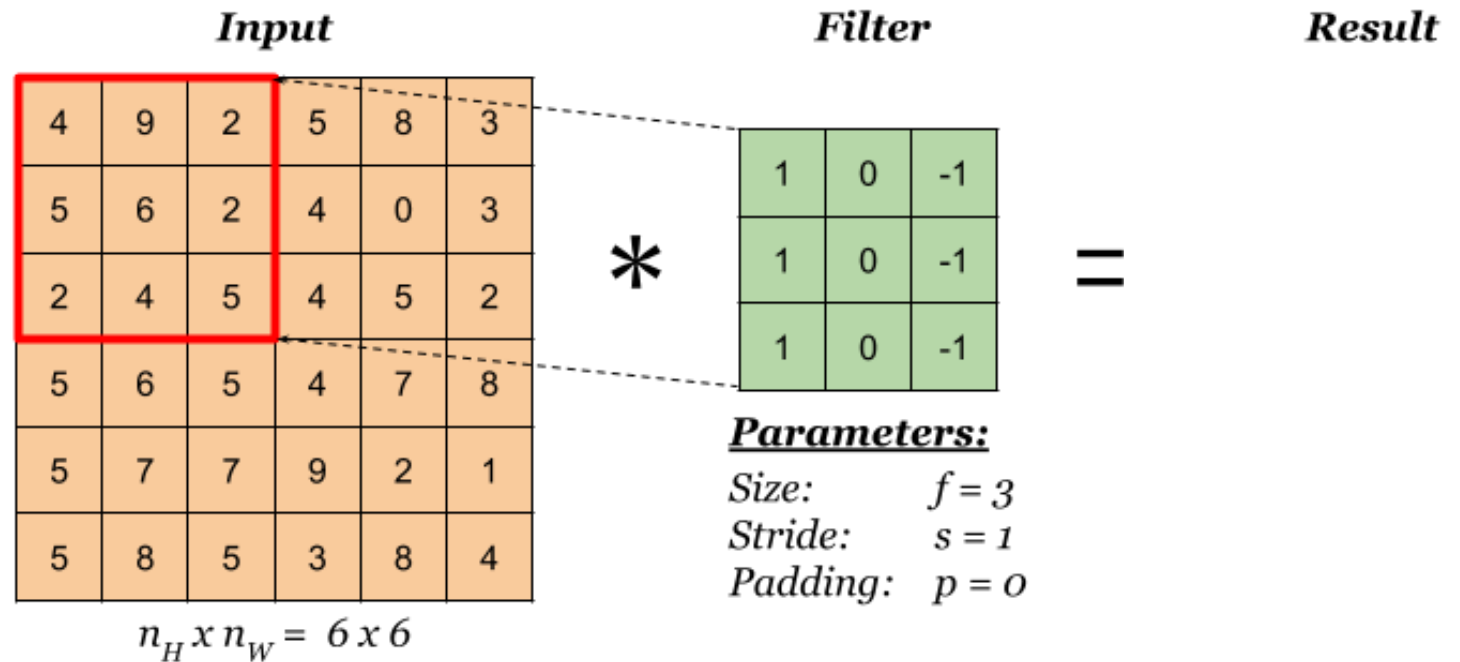
s : convolution stride size

$$\begin{aligned} \bullet \quad n_{out} &= \left\lfloor \frac{6 - 3 + 2 \cdot 1}{2} \right\rfloor + 1 = \left\lfloor \frac{5}{2} \right\rfloor + 1 \\ &= 2 + 1 = 3 \end{aligned}$$



Determine the Output Dimension

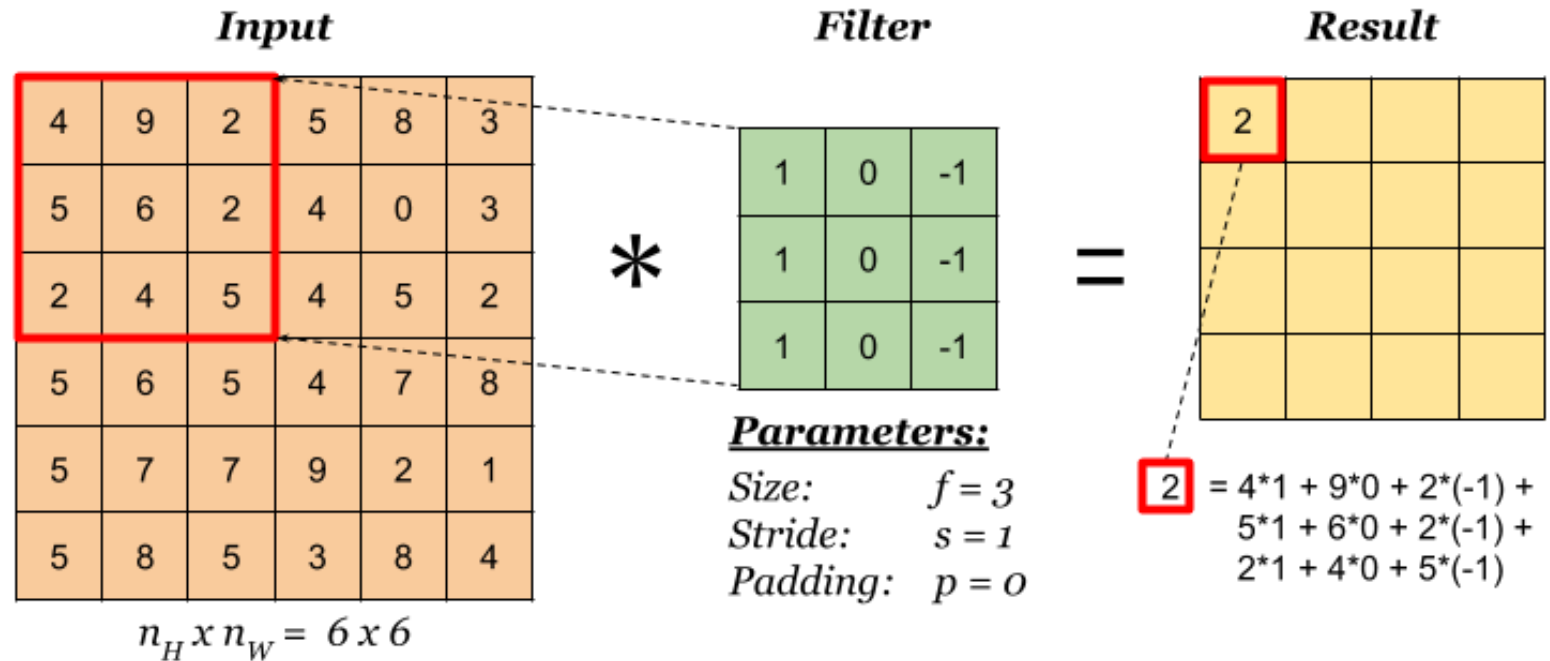
• $n_{out} = ?$





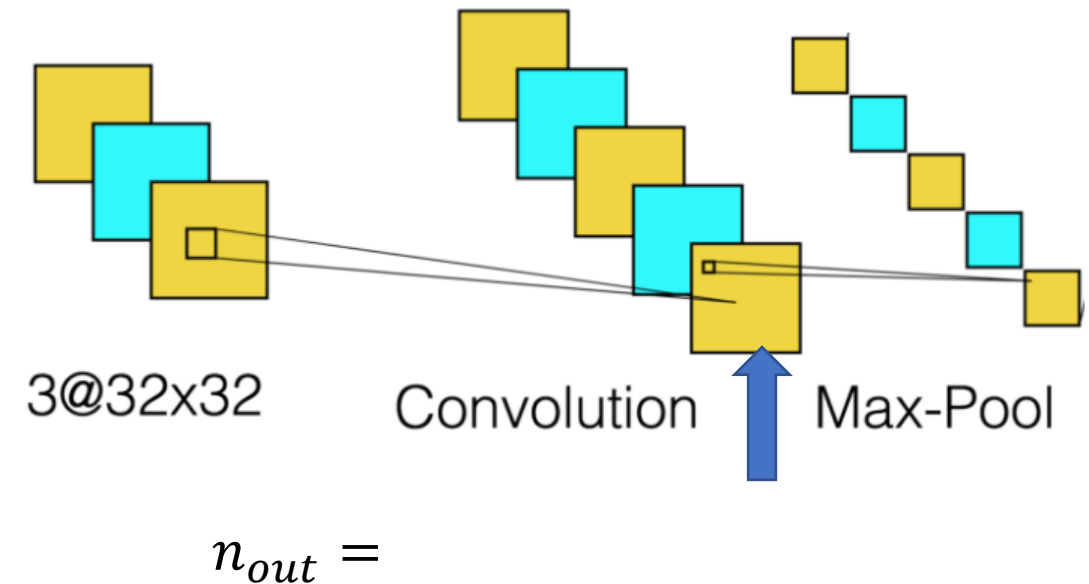
Determine the Output Dimension

$$\begin{aligned}
 \bullet \ n_{out} &= \left\lfloor \frac{6-3+2*0}{1} \right\rfloor + 1 \\
 &= \left\lfloor \frac{3}{1} \right\rfloor + 1 \\
 &= 4
 \end{aligned}$$




Determine the Output Dimension

- Input image, I with dimensions $(32 \times 32 \times 3)$
- Convolution Layer
 - A filter size 3×3
 - Stride is 1
 - Valid padding, and
 - Depth/feature maps are 5 ($D = 5$)
- **Output dimensions = ?**



Determine the Output Dimension

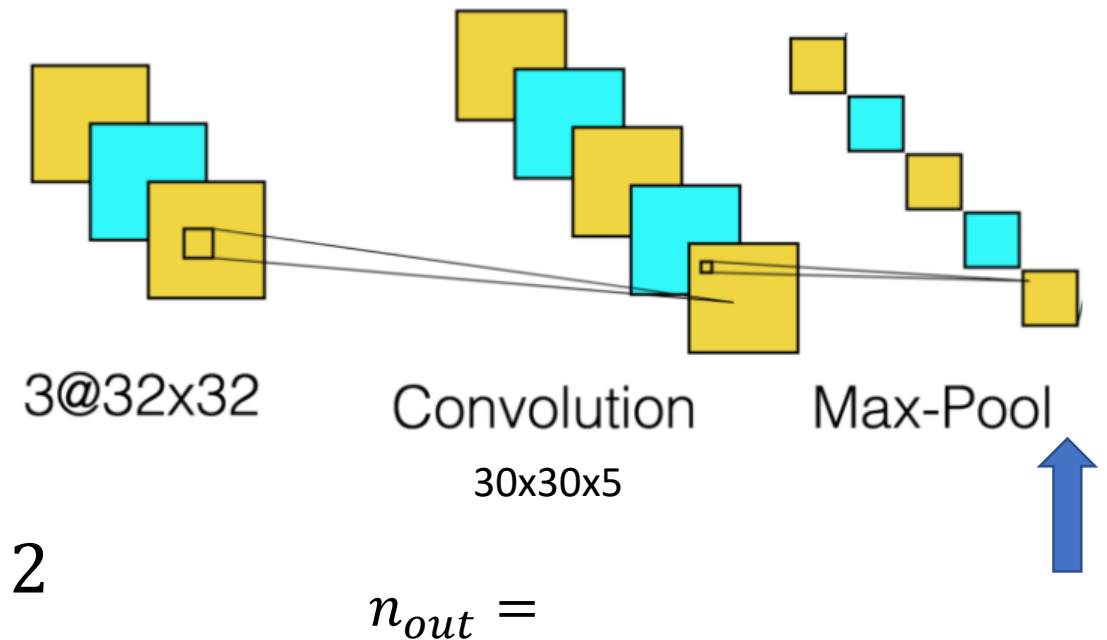
- Input image, I with dimensions $(32 \times 32 \times 3)$
- Convolution Layer
 - A filter size 3×3
 - Stride is 1 ($s=1$)
 - Valid padding ($p=0$), and
 - Depth/feature maps are 5 ($D=5$)
- **Output dimensions = $30 \times 30 \times 5$**
- **After Pooling?**



$$n_{out} = \left\lfloor \frac{32 - 3 + 2 * 0}{1} \right\rfloor + 1 = \left\lfloor \frac{29}{1} \right\rfloor + 1 = 30$$

Determine the Output Dimension

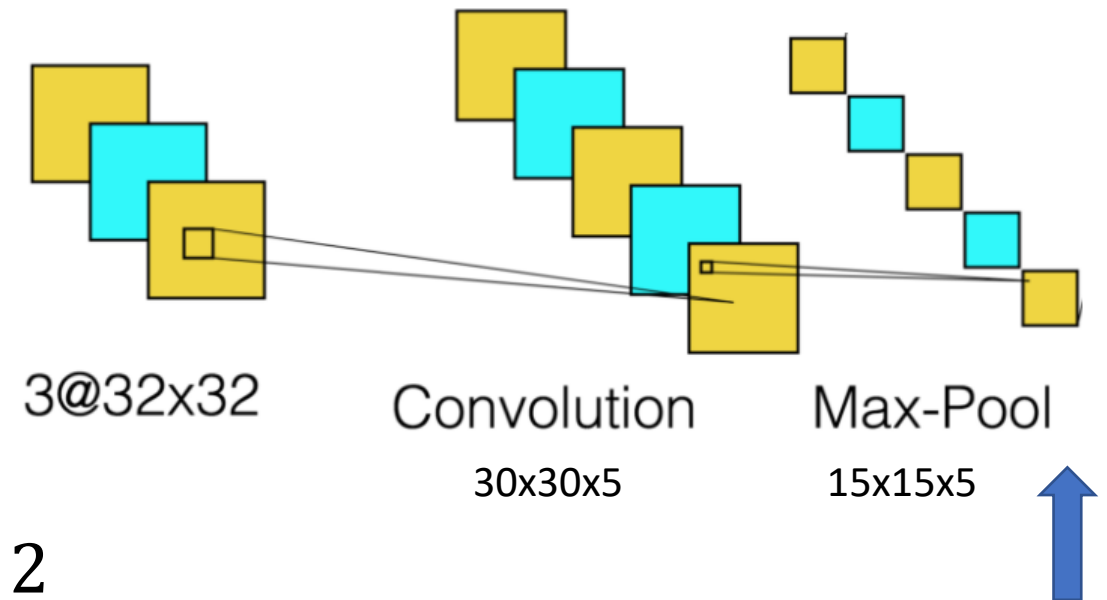
- Input to Pooling Layer (30x30x5)
- After Pooling with
 - Filter size $k \times k$
 - Stride s
- $$n_{out} = \left\lfloor \frac{n_{in} - k}{s} \right\rfloor + 1$$
- **Eg,** Pooling with, Filter size 2×2 , Stride 2
- Output dimensions =





Determine the Output Dimension

- Input to Pooling Layer (30x30x5)
- After Pooling with
 - Filter size $k \times k$
 - Stride s
- $n_{out} = \left\lfloor \frac{n_{in} - k}{s} \right\rfloor + 1$
- **Eg**, Pooling with, Filter size 2×2 , Stride 2
- Output dimensions = 15x15x5



$$n_{out} = \left\lfloor \frac{30 - 2}{2} \right\rfloor + 1 = \left\lfloor \frac{28}{2} \right\rfloor + 1 = 15$$



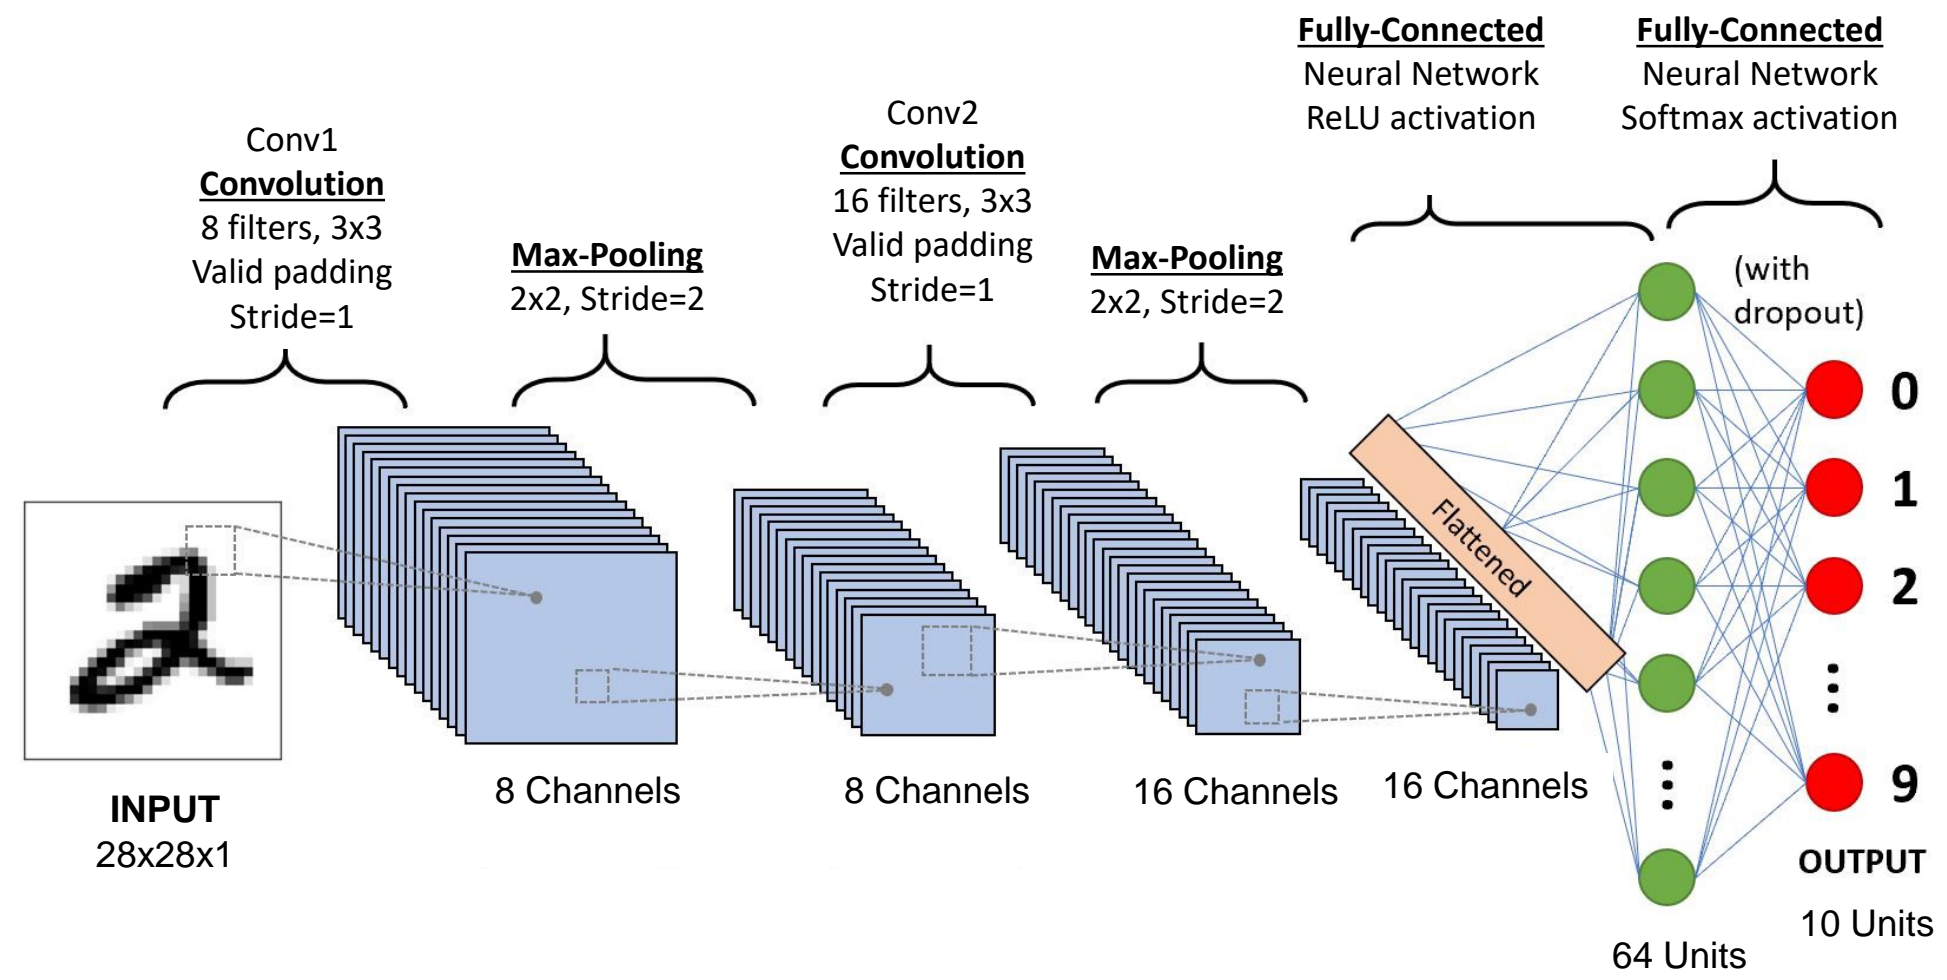
Typical CNN Model

• Conv1

$$n_{out} = \left\lfloor \frac{n_{in} - k + 2 * p}{s} \right\rfloor + 1$$

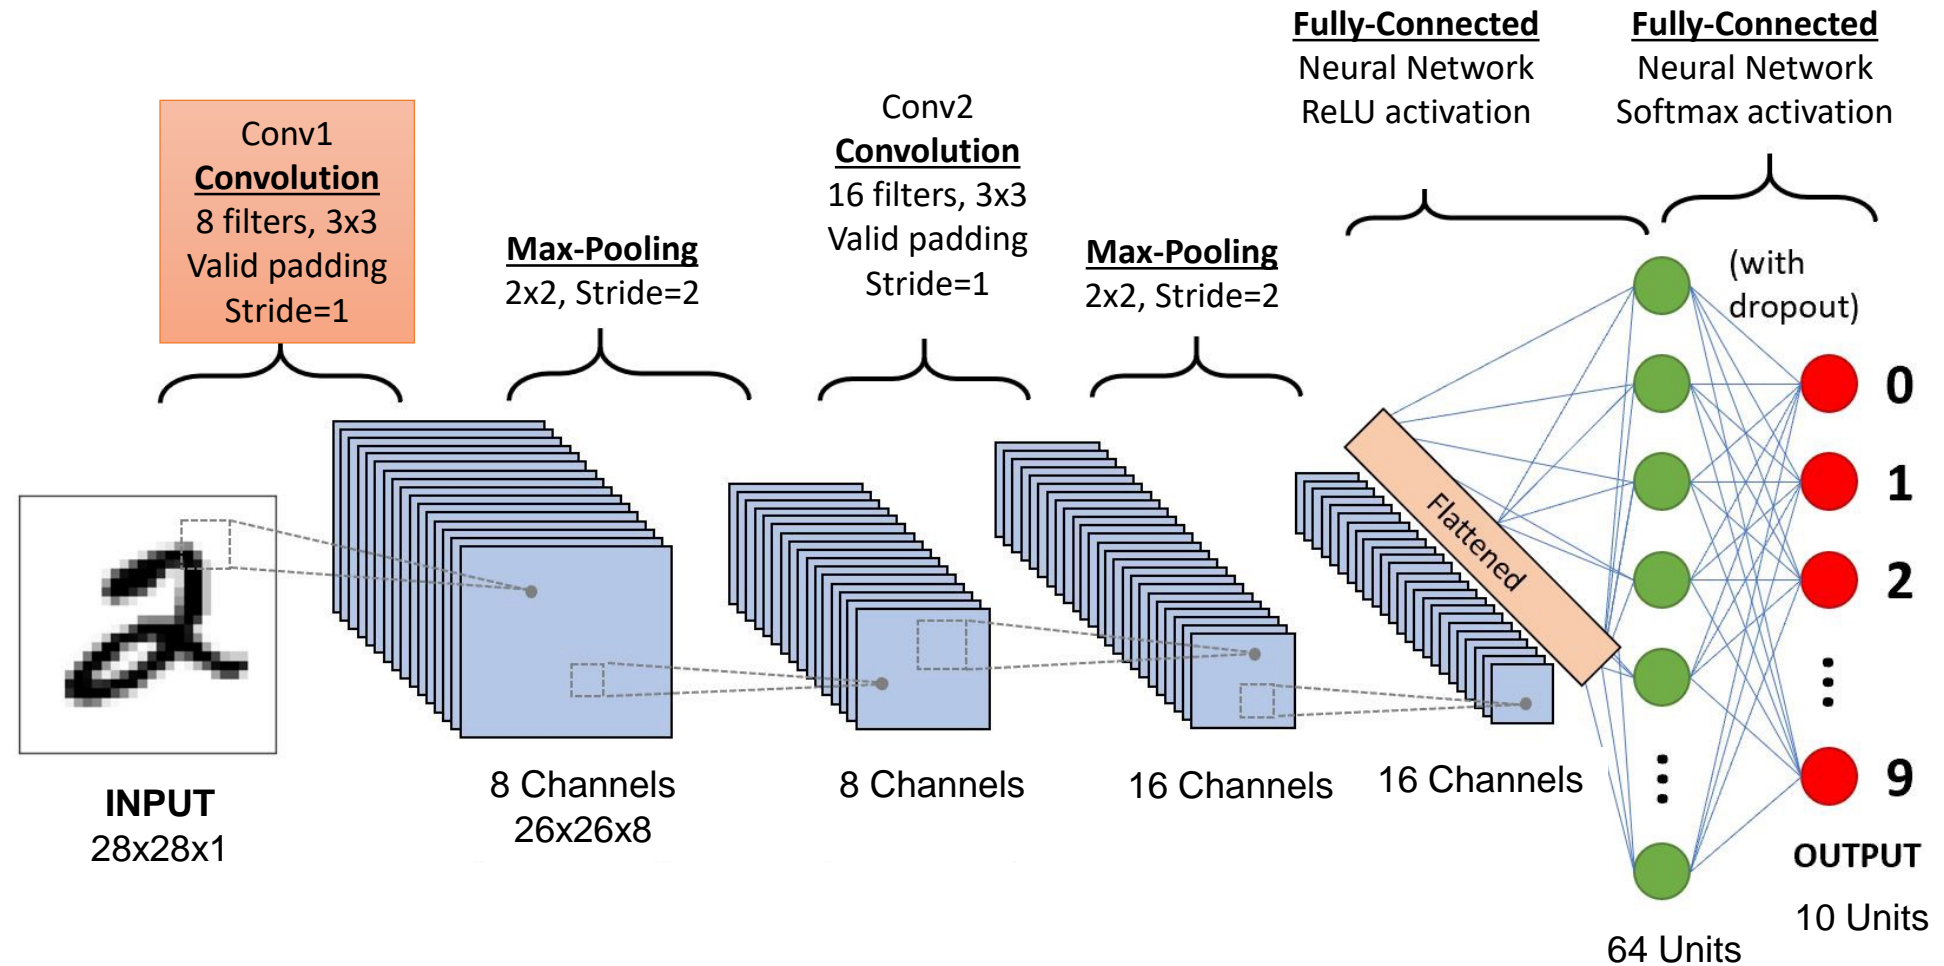
=?

• o/p:



Typical CNN Model

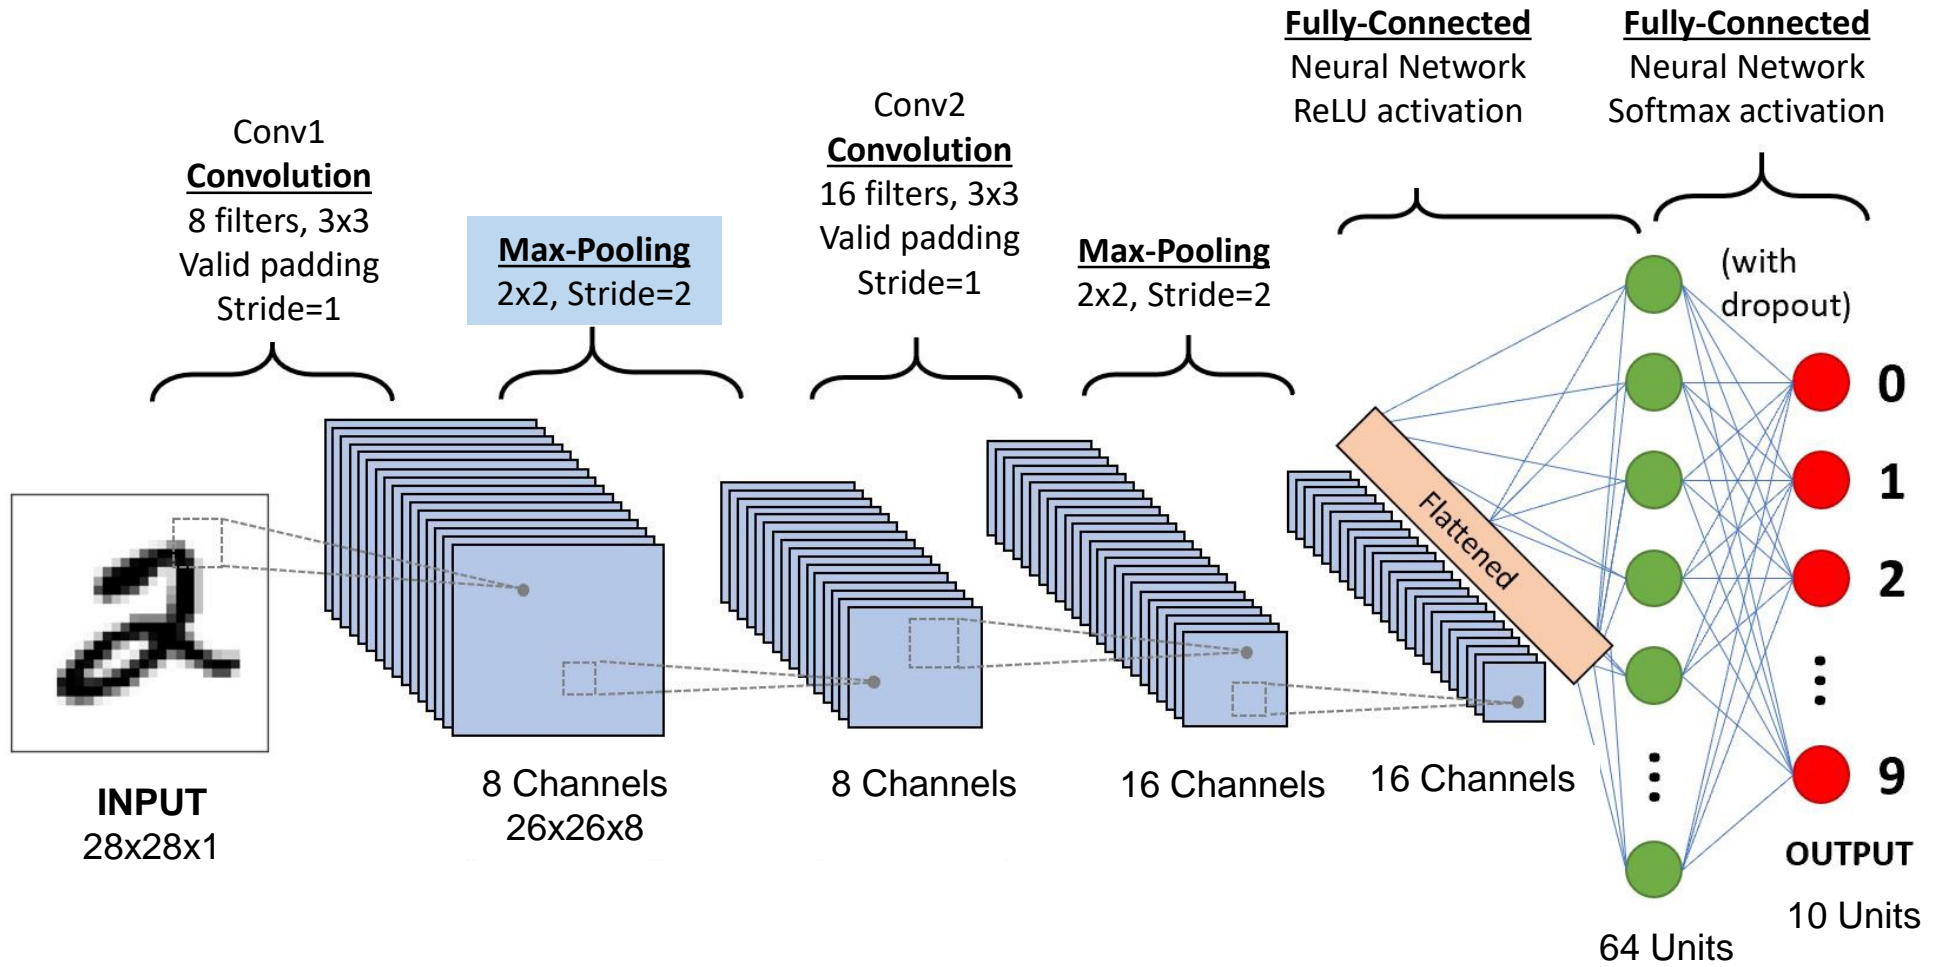
- $\text{Conv1} \left\lfloor \frac{28-3+2*0}{1} \right\rfloor + 1 = 26$
- o/p: $26 \times 26 \times 8$
- Param: $3 \times 3 \times 1 \times 8 + 8 = 80$
- 3x3 filter for 1 channel, 8 such filters and 8 biases





Typical CNN Model

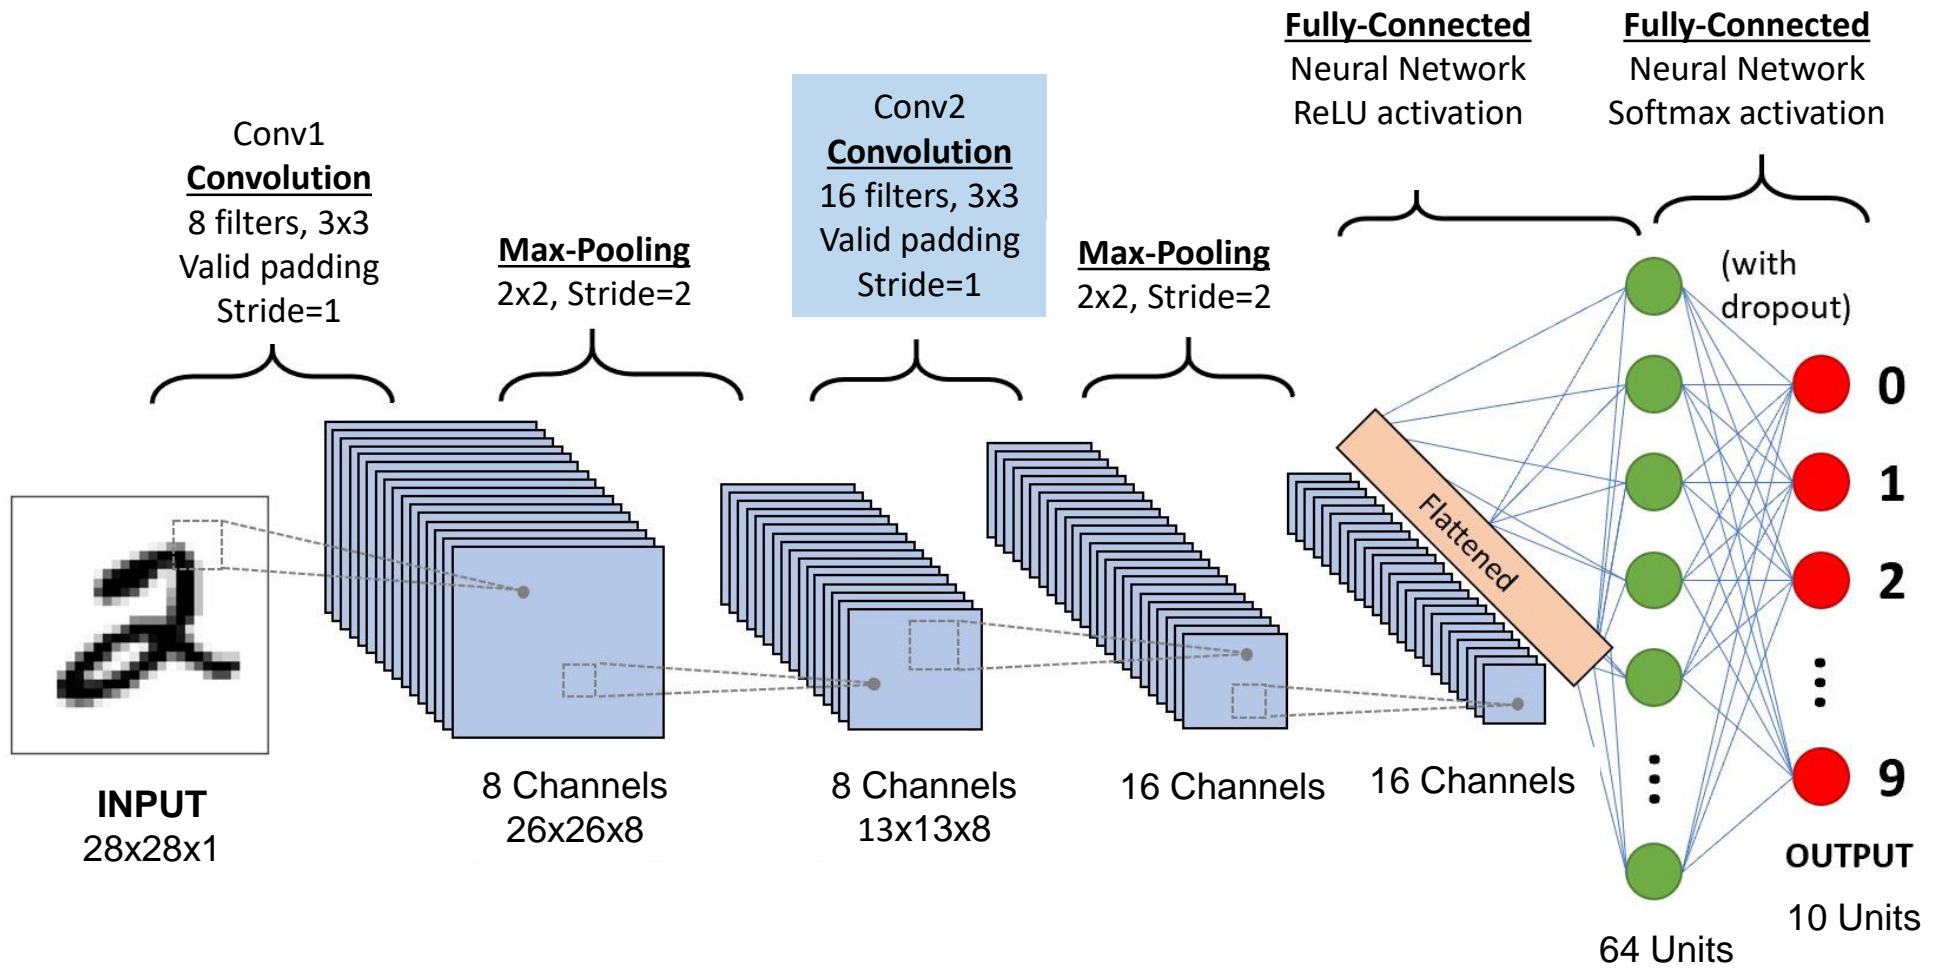
- Conv1:
 $26 \times 26 \times 8$
- Max-Pool
 $\left\lfloor \frac{n_{in} - k}{s} \right\rfloor + 1 = ?$
o/p: ?





Typical CNN Model

- Conv1:
 $26 \times 26 \times 8$
- Max-Pool
 $\left\lfloor \frac{26-2}{2} \right\rfloor + 1 = 13$
o/p: $13 \times 13 \times 8$
- Conv2, $n_{out} =$
 $\left\lfloor \frac{n_{in}-k+2*p}{s} \right\rfloor + 1$

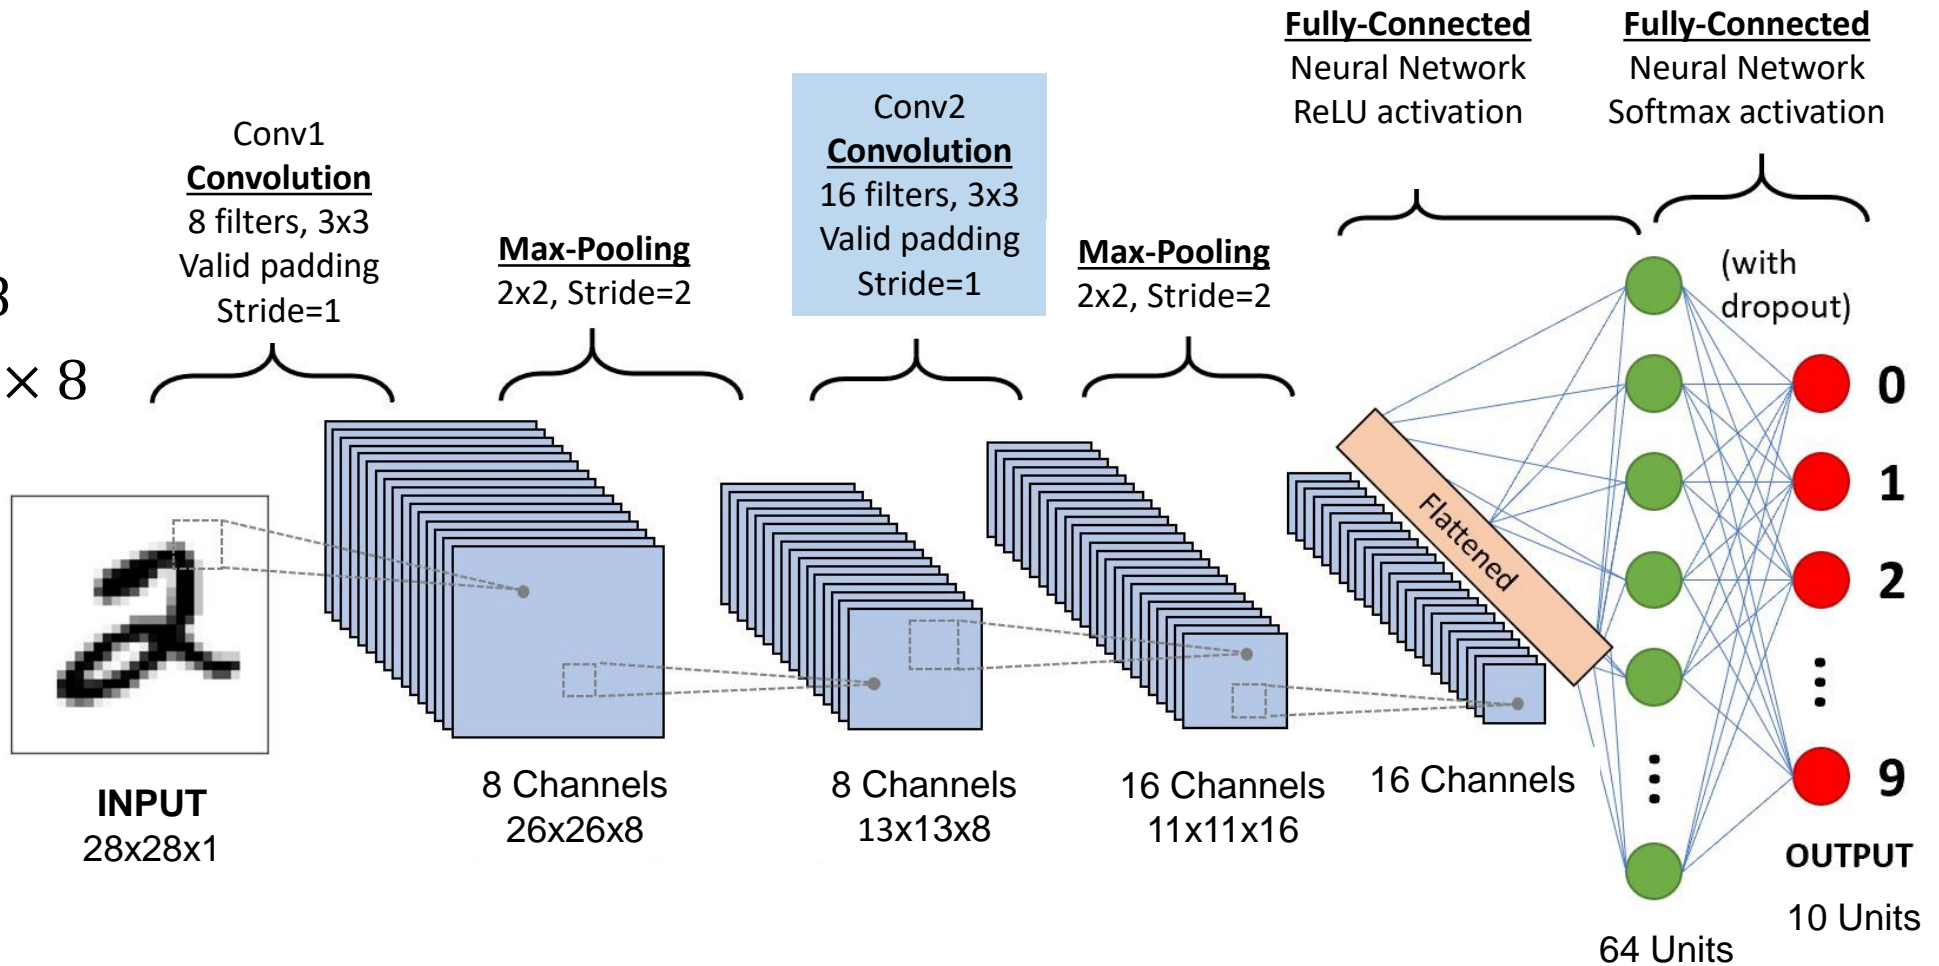




Typical CNN Model

- Conv1: $26 \times 26 \times 8$
- Max-Pool: $13 \times 13 \times 8$
- Conv2:

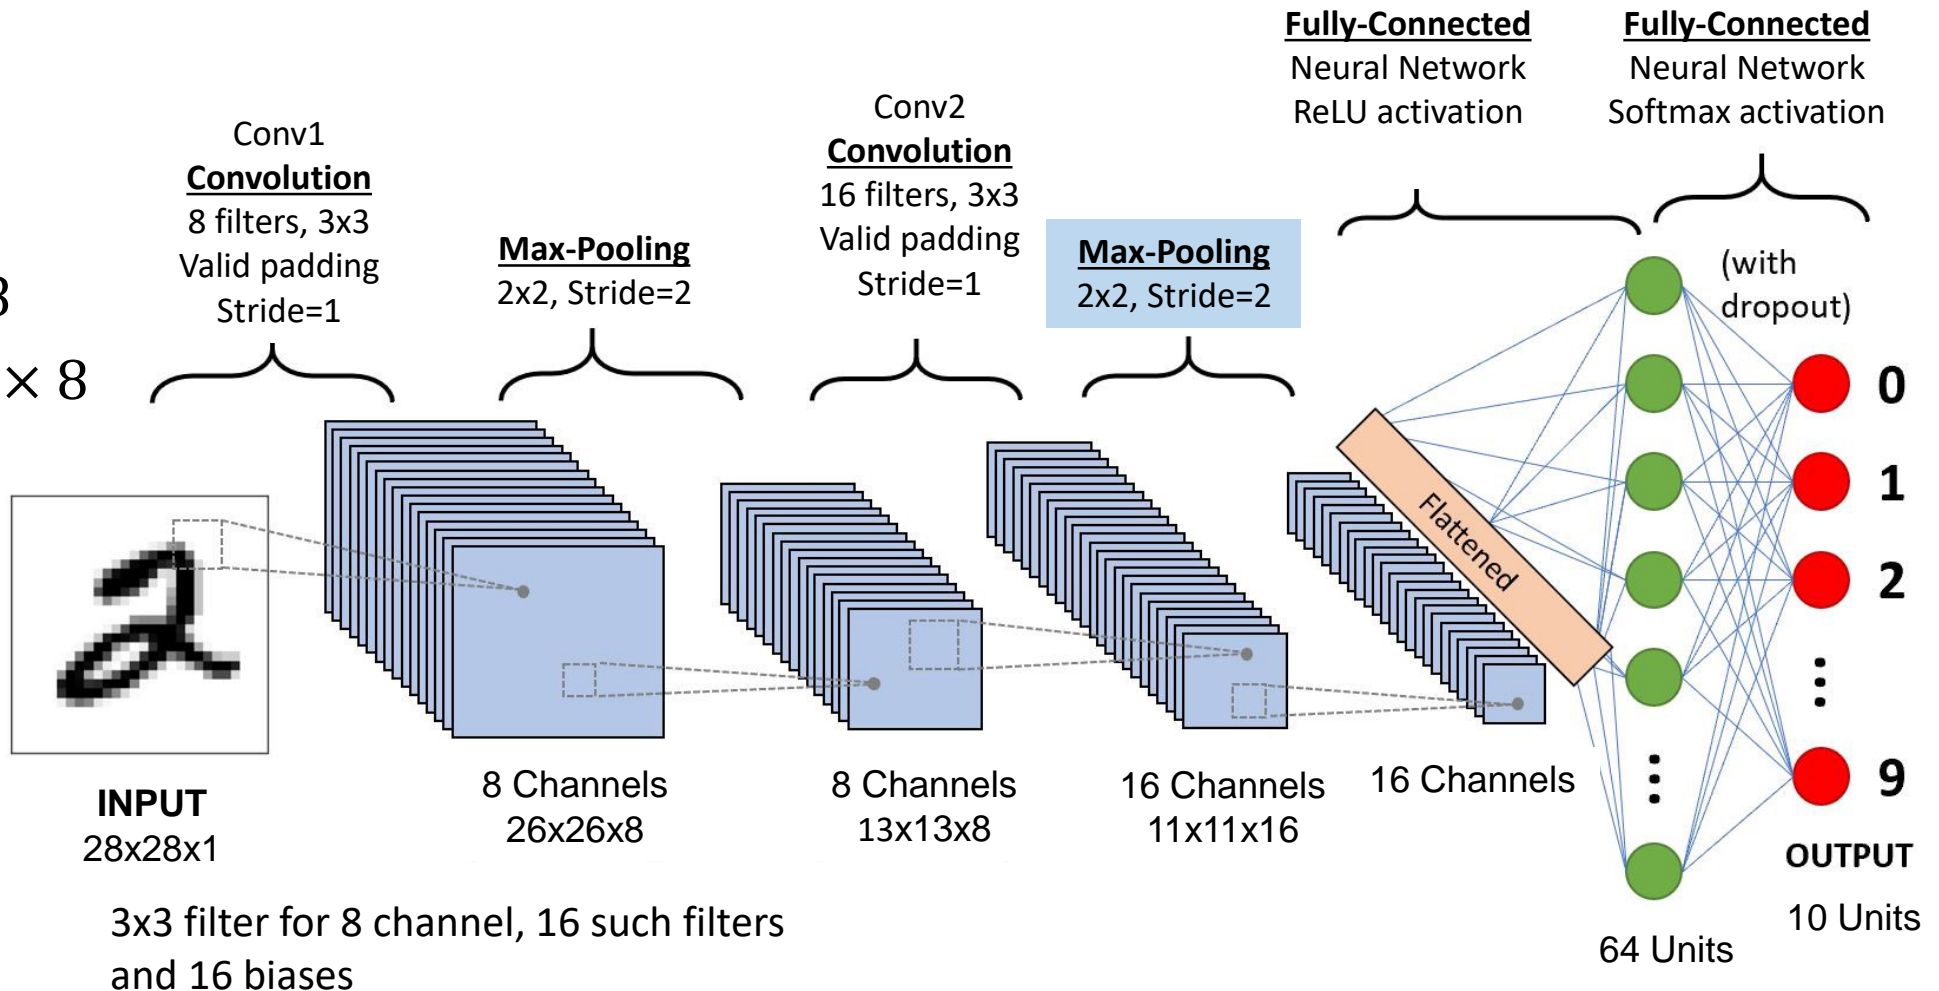
$$\left\lfloor \frac{13-3+2*0}{1} \right\rfloor + 1 = 11$$
- o/p: $11 \times 11 \times 16$
- Param?





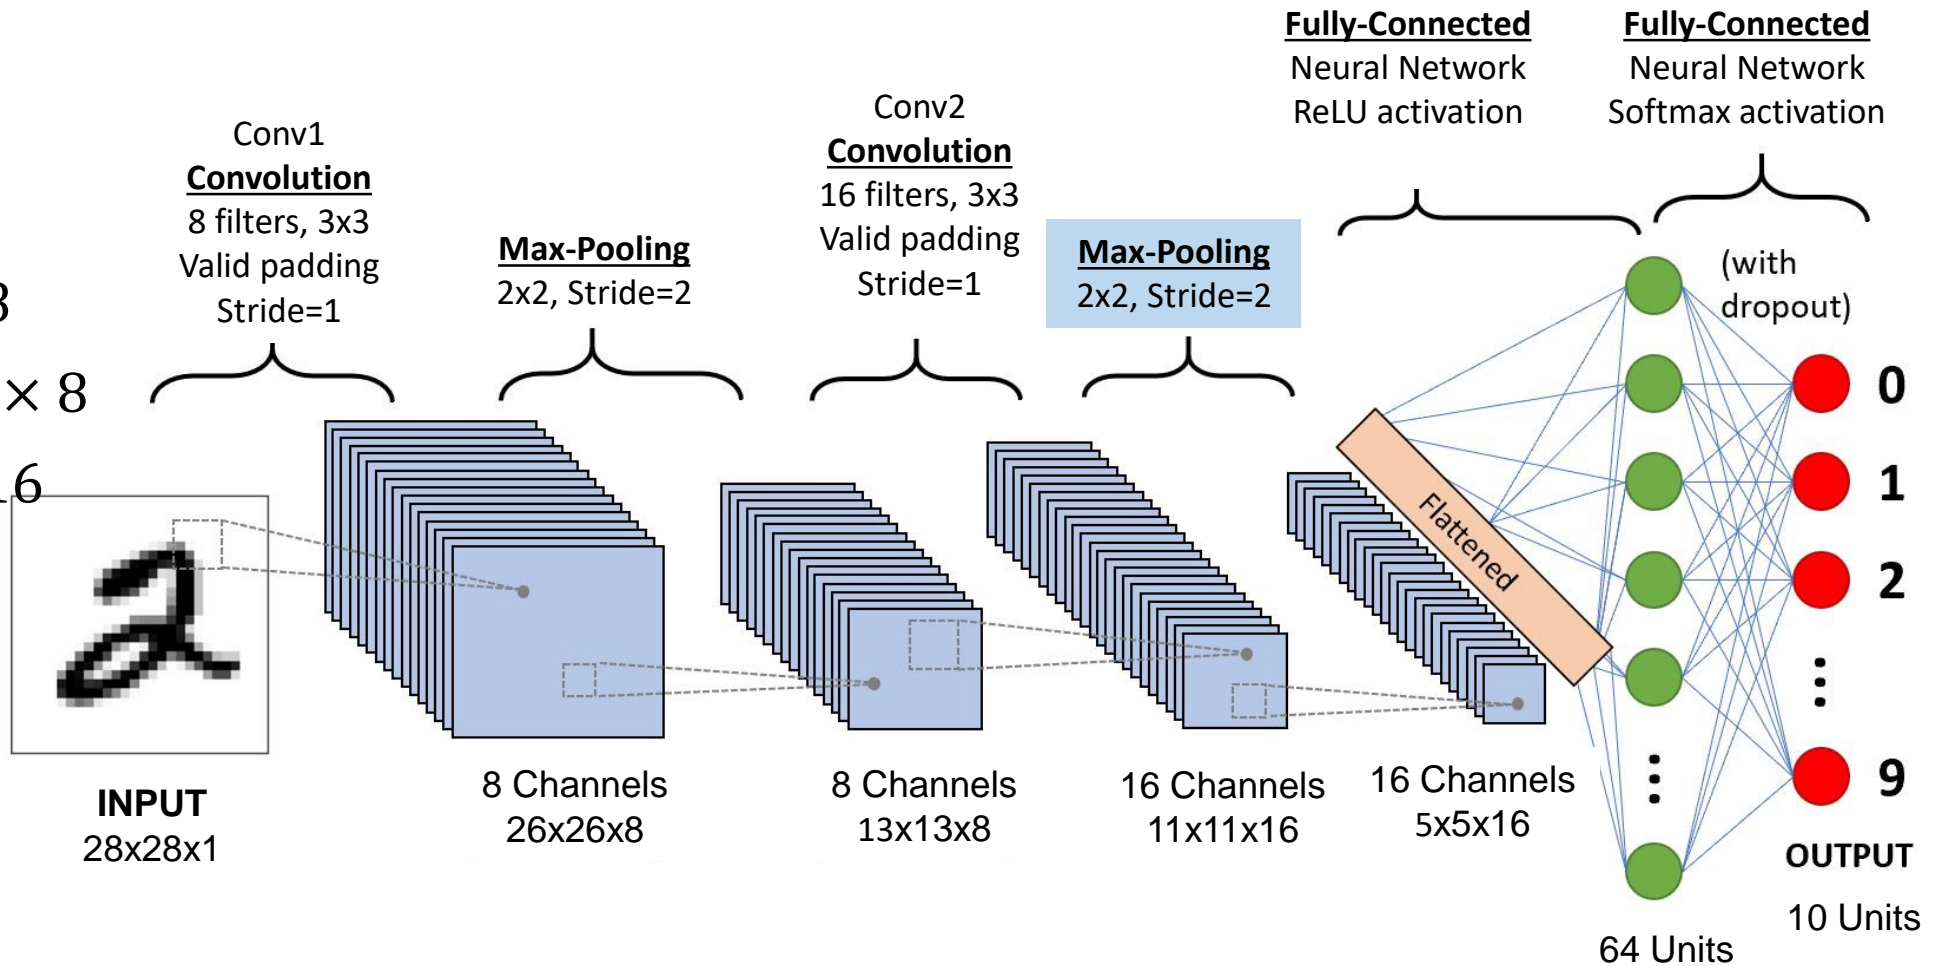
Typical CNN Model

- Conv1: $26 \times 26 \times 8$
 - Max-Pool: $13 \times 13 \times 8$
 - Conv2:
- $$\left\lfloor \frac{13-3+2*0}{1} \right\rfloor + 1 = 11$$
- o/p: $11 \times 11 \times 16$
- Param=
- $$3 \times 3 \times 8 \times 16 + 16 = 1168$$



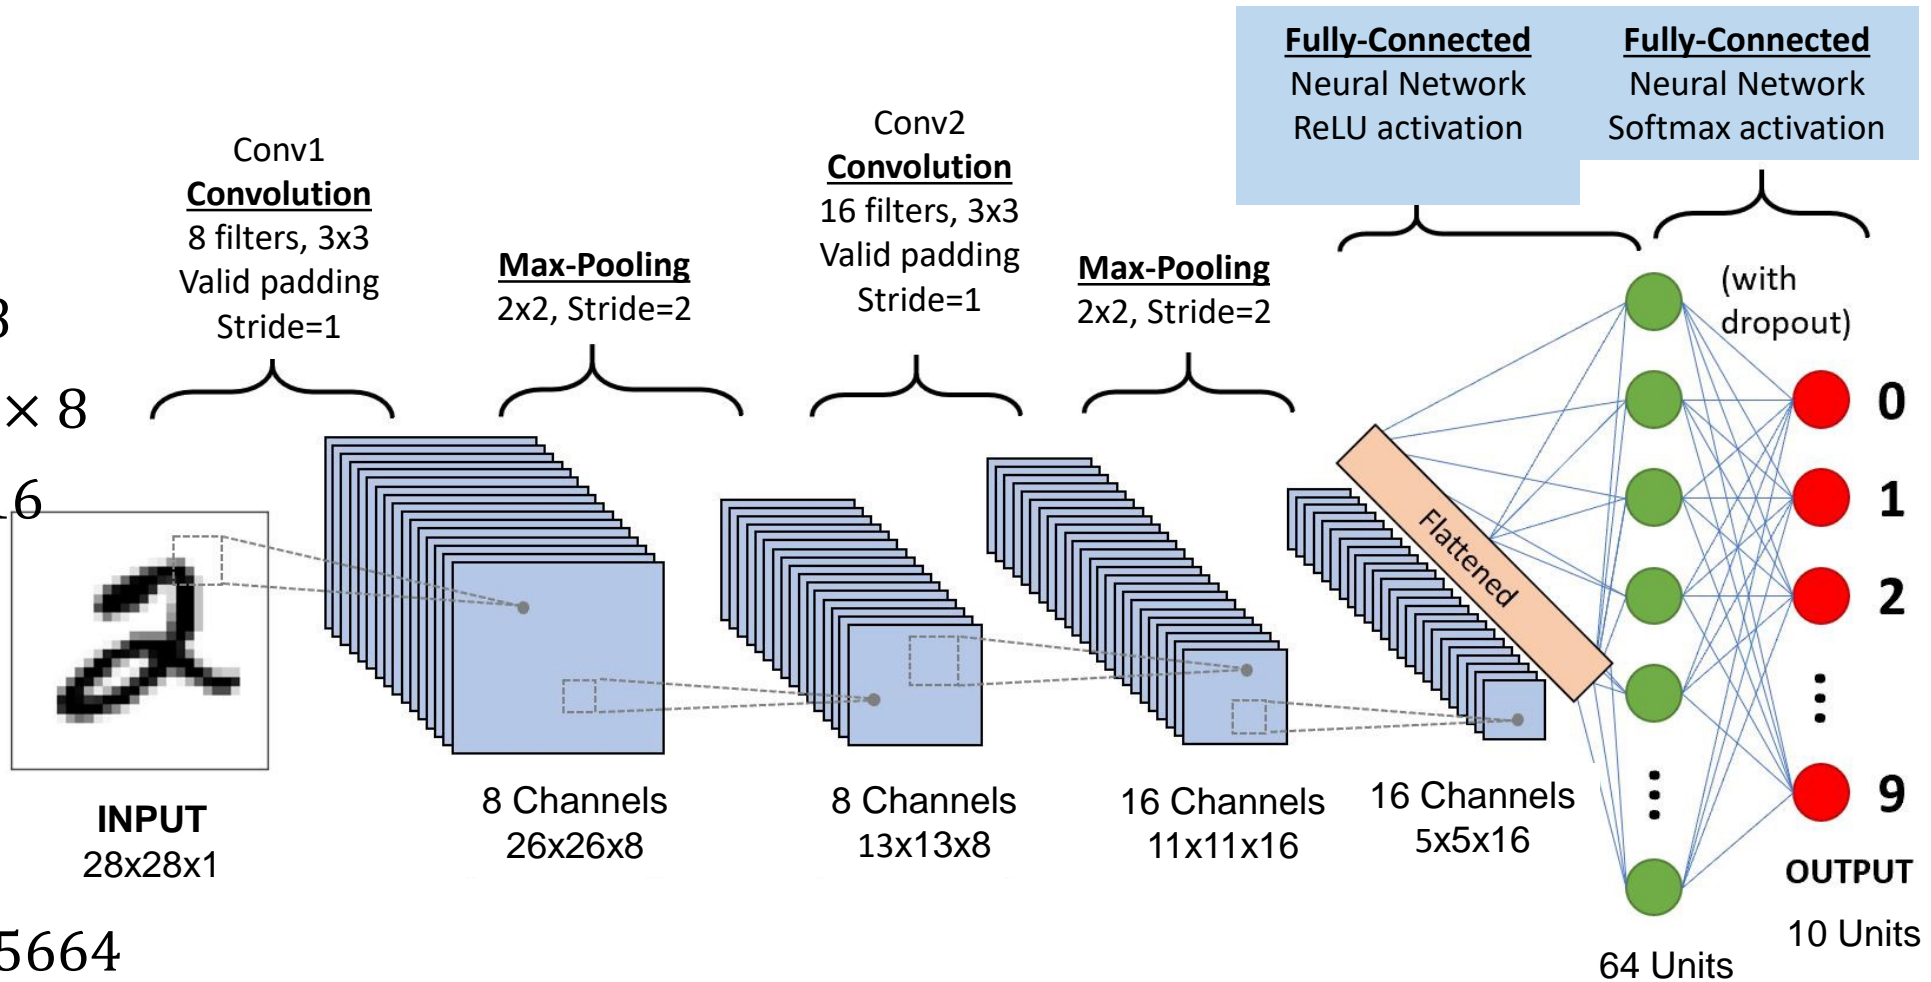
Typical CNN Model

- Conv1: $26 \times 26 \times 8$
- Max-Pool: $13 \times 13 \times 8$
- Conv2: $11 \times 11 \times 16$
- Max-Pool: $5 \times 5 \times 16$
- $\left\lfloor \frac{11-2}{2} \right\rfloor + 1 = 5$



Typical CNN Model

- Conv1: $26 \times 26 \times 8$
- Max-Pool: $13 \times 13 \times 8$
- Conv2: $11 \times 11 \times 16$
- Max-Pool:
 $5 \times 5 \times 16$
- $5 \times 5 \times 16 = 400$
- FC1:
 $(400 + 1) \times 64 = 25664$
- FC2: $(64+1) \times 10 = 650$





Parameters and Hyperparameters

	Parameters	Hyperparameters
Convolution layer	Kernels	Kernel size, number of kernels, stride, padding, activation function
Pooling layer	None	Pooling method, filter size, stride, padding
Fully connected layer	Weights	Number of weights, activation function
Others		Model architecture, optimizer, learning rate, loss function, mini-batch size, epochs, regularization, weight initialization, dataset splitting

Note that a parameter is a variable that is automatically optimized during the training process and a hyperparameter is a variable that needs to be set beforehand






CNN Playground

Click the link below or copy paste the URL in your browser

<https://poloclub.github.io/cnn-explainer/>

With the CNN Explainer you can Learn and implement Convolutional Neural Network (CNN) in your browser! With real sample image dataset

Understanding Receptive field




- The human visual system consists of millions of neurons, where each one captures different information.
- Defined as neuron's receptive field as the patch of the total field of view Or what information a single neuron has access to..

ImageSource: <https://www.brainhq.com/brain-resources/brain-connection>



Receptive Field in Deep Learning




5x5

3x3

Effective Receptive Field

- Defined as a size of region in input that produces features. Basically, it is a measure of association of an output feature (of any layer) to the input region (patch).
- The idea of receptive fields applies to local operations (i.e. convolution, pooling).
- A convolutional unit only depends on a local region (patch) of the input.
- That's why RF never referred on fully connected layers since each unit has access to all the input region.


Why do we care for Receptive Field?



The green and the orange one. Which one would you like to have in your architecture?

Image Source: <https://developer.nvidia.com/blog/image-segmentation-using-digits-5/>

Why do we care for Receptive Field?




Therefore, our goal is to design a convolutional model so that we ensure that its RF covers the entire relevant input image region.

Image Source: <https://developer.nvidia.com/blog/image-segmentation-using-digits-5/>

How to increase receptive field in a convolutional network?

- Add more convolutional layers (make the network deeper)
- Add pooling layers or higher stride convolutions (sub-sampling)
- Use dilated convolutions



It is a technique that expands the kernel (input) by inserting holes between its consecutive elements.





Parameter Sharing




- Parameter sharing refers to using the same parameter for more than one function in a model
- Kernel is reused (by sliding) when calculating the layer o/p
- Less weights to store & train





Equivalent Representation

Representation



Image

- Parameter sharing causes the layer to have a property called **equivariance to translation**
- Convolution creates 2-D map of where certain features appear in the input
- If we move the object in the input, its representation will move the same amount in the output
- Eg: Same kernel for Edge Detection wherever the edge occurs in the image




1D, 2D Convolution

- 2D Convolution
 - 2-directions (x,y) to calculate conv
 - input = $(W \times H \times c)$, d filters $(k \times k \times c)$ output = $(W_1 \times H_1 \times d)$
 - Eg: Image data (gray or color)
- 1D Convolution
 - 1-direction (time) to calculate conv
 - input = (time-step $\times c$),

d filters $(k \times c)$, output (time-step1 $\times d$)

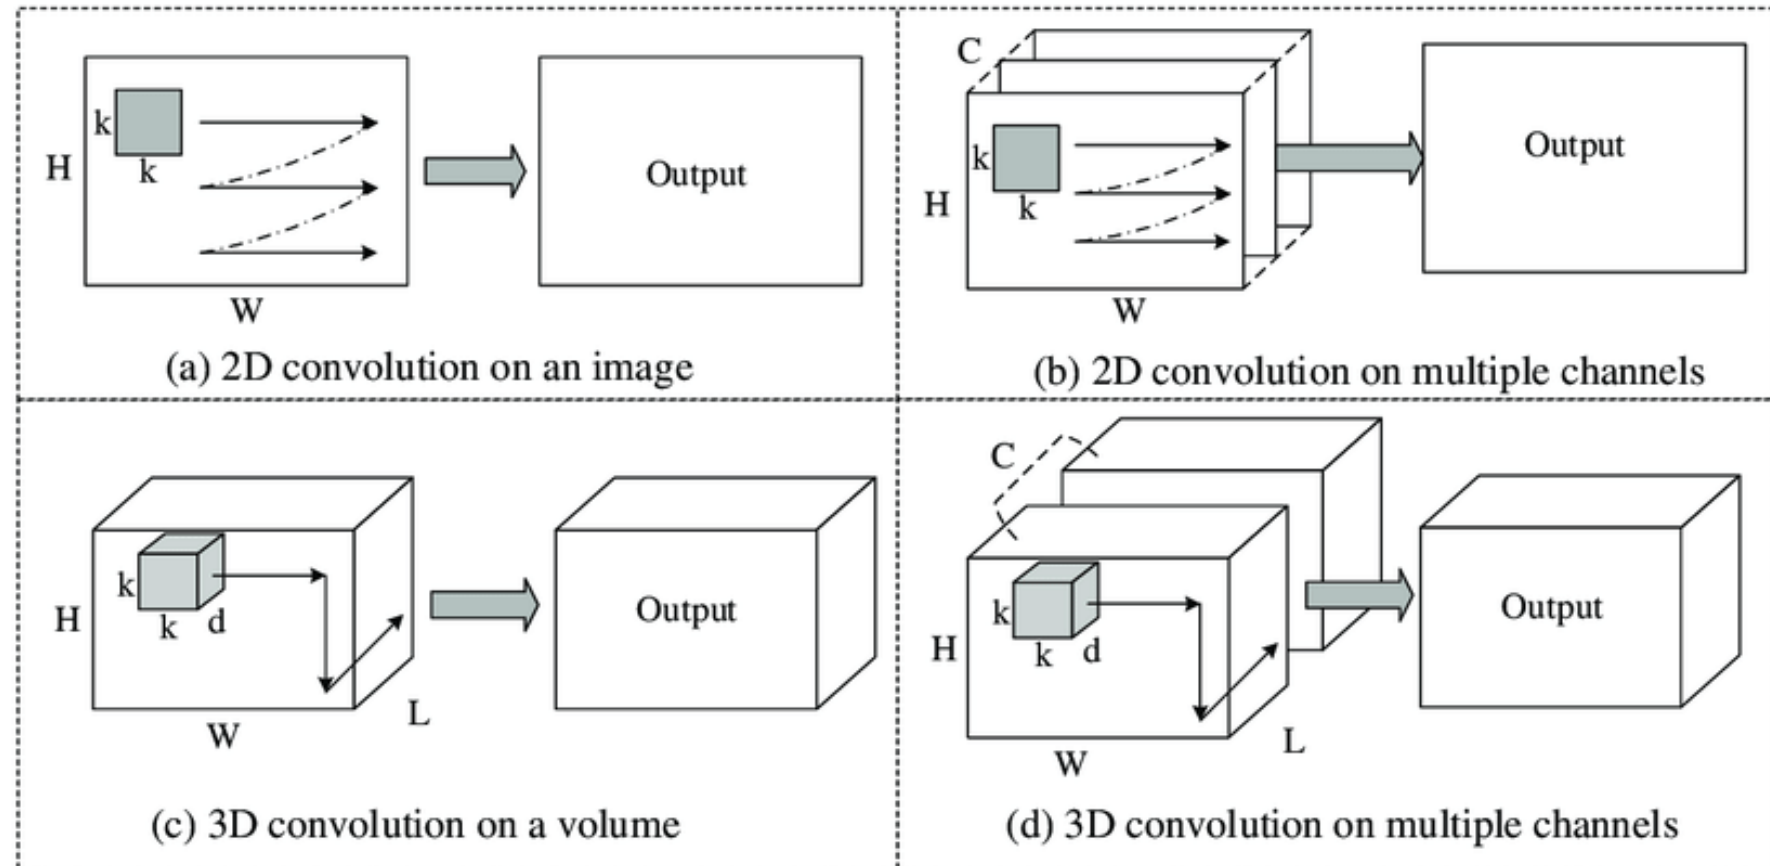
Eg: Time-series data, text analysis





2D, 3D Convolution

- 3D Convolution
- 3-directions (x,y,z) to calculate conv
- input ($W \times H \times L \times C$), m filters ($k \times k \times d$) output ($W_1 \times H_1 \times L_1 \times m$)
- Eg: MRI data, Videos





1 D CNN

Kernel size = 3



Input shape = 2D


Batch size = None

Width = Time axis = 7

Feature map/ channels = 1

2 D CNN

Kernel size = (3,3)



Input shape = 3D


Height=5

Width = 7

Feature map/ channels = 1

3 D CNN

Kernel size = (3,3,2)



Input shape = 4D

Height=6

Width = 6

Feature map/ channels
=depth=1



1 D CNN

Kernel size = 3



Input shape = 2D


Batch size = None

Width = Time axis = 7

Feature map/ channels = 1

2 D CNN

Kernel size = (3,3)



Input shape = 3D


Height=5

Width = 7

Feature map/ channels = 1

3 D CNN

Kernel size = (3,3,2)



Input shape = 4D

Height=6

Width = 6

Feature map/ channels = depth=1



1 D CNN

Kernel size = 3



Input shape = 2D


Batch size = None

Width = Time axis = 7

Feature map/ channels = 1

2 D CNN

Kernel size = (3,3)



Input shape = 3D


Height=5

Width = 7

Feature map/ channels = 1

3 D CNN

Kernel size = (3,3,2)



Input shape = 4D

Height=6

Width = 6

Feature map/ channels = depth=1